



THE UNIVERSITY OF QUEENSLAND
A U S T R A L I A

OPTIMIZING PERFORMANCE
IN GAUSSIAN PROCESSES

MICHAEL CICCOTOSTO-CAMP

SUPERVISOR: FRED (FARBOD) ROOSTA
Co-SUPERVISORS: ANDRIES POTGIETER
YAN ZHAO

BACHELOR OF MATHEMATICS (HONOURS)
JUNE 2022

THE UNIVERSITY OF QUEENSLAND
SCHOOL OF MATHEMATICS AND PHYSICS

CONTENTS

ACKNOWLEDGEMENTS	iii
SYMBOLS AND NOTATION	iv
INTRODUCTION	1
GAUSSIAN PROCESSES	4
0.1. KERNELS	4
0.2. REPRODUCING KERNEL HILBERT SPACES	5
0.3. GAUSSIAN RADIAL BASIS KERNEL	7
0.4. KERNEL MACHINES	7
0.4.1. INTRODUCTION TO SUPPORT VECTOR MACHINES FOR BINARY CLASSIFICATION	8
0.4.2. INTRODUCTION TO GAUSSIAN PROCESSES FOR REGRESSION	9
0.5. GAUSSIAN PROCESSES FOR REGRESSION	13
0.5.1. NOISE-FREE OBSERVATIONS	15
0.5.2. PREDICTION WITH NOISY OBSERVATIONS	16
0.6. GAUSSIAN PROCESSES FOR CLASSIFICATION	17
0.6.1. LINEAR MODELS FOR CLASSIFICATION	18
0.6.2. LAPACE APPROXIMATION FOR POSTERIOR	20
0.6.3. PREDICTIONS	22
1. APPLICATIONS AND RESULTS	24
1.1. GAUSSIAN PROCESSES PREDICTION REVIEWED	24
1.2. EXPERIMENTAL SETUP	24
REFERENCES	25
APPENDIX A. SUPPLEMENTARY RESULTS	30
A.1. GRAM MATRIX SPECTRAL VALUES	30
A.2. NYSTROM SCORES	32
A.3. RIDGE LEVERAGE SCORES	35
A.4. NYSTROM ERRORS	38

ACKNOWLEDGEMENTS

I would like to deeply thank my supervisor Dr. Masoud Kamgarpour for his advice and all of his time spent with me. I consider myself lucky and am glad to have been his student for my honours year. I would also like to thank my co-supervisor Dr. Anna Puskás for the same reasons. A special thanks to Dr. Valentin Buciumas for his time spent teaching me while he was at The University of Queensland.

SYMBOLS AND NOTATION

Matrices are capitalized bold face letters while vectors are lowercase bold face letters.

<i>Syntax</i>	<i>Meaning</i>
\triangleq	An equality which acts as a statement
$ A $	The determinate of a matrix.
$\langle \cdot, \cdot \rangle_{\mathcal{H}}$	The inner product with respect to the Hilbert space \mathcal{H} , sometimes abbreviated as $\langle \cdot, \cdot \rangle$ if the Hilbert space is clear from context.
$\ \cdot\ _{\mathcal{V}}$	The norm of a vector with respect to the vector space \mathcal{V} , sometimes abbreviated as $\ \cdot\ $ if the vector space is clear from context.
x^T, X^T	The transpose operator.
x^*, X^*	The hermitian operator.
$a.*b$ or $A.*B$	Element-wise vector (matrix) multiplication, similar to Matlab.
\propto	Proportional to.
∇ or ∇_f	The partial derivative (with respect to f).
∇	The Hessian.
\sim	Distributed according to, example $x \sim \mathcal{N}(0, 1)$
0 or 0_n or $0_{n \times m}$	The zero vector (matrix) of appropriate length (size) or the zero vector of length n or the zero matrix with dimensions $n \times m$.
1 or 1_n or $1_{n \times m}$	The one vector (matrix) of appropriate length (size) or the one vector of length n or the one matrix with dimensions $n \times m$.
$\mathbb{1}_{n \times m}$	The matrix with ones along the diagonal and zeros on off diagonal elements.

$\mathbf{A}_{(.,.)}$	Index slicing to extract a submatrix from the elements of $\mathbf{A} \in \mathbb{R}^{n \times m}$, similar to indexing slicing from the python and Matlab programming languages. Each parameter can receive a single value or a 'slice' consisting of a start and an end value separated by a semicolon. The first and second parameter describe what row and columns should be selected, respectively. A single value means that only values from the single specified row/column should be selected. A slice tells us that all rows/columns between the provided range should be selected. Additionally if now start and end values are specified in the slice then all rows/columns should be selected. For example, the slice $\mathbf{A}_{(1:3,j:j')}$ is the submatrix $\mathbb{R}^{3 \times (j'-j+1)}$ matrix containing the first three rows of \mathbf{A} and columns j to j' . As another example, $\mathbf{A}_{(:,j)}$ is the j^{th} column of \mathbf{A} .
\mathbf{A}^\dagger	Denotes the unique psuedo inverse or Moore-Penore inverse of \mathbf{A} .
\mathbb{C}	The complex numbers.
C	The classes in a classification problem.
$\text{cholesky}(\mathbf{A})$	A function to compute the Cholesky decomposition of the matrix \mathbf{A} , where $\mathbf{L}\mathbf{L}^\top = \mathbf{A}$.
$\text{cov}(\mathbf{f})$	Gaussian process posterior covariance.
d	The number of features in the data set.
D	The dimension of the feature space of the feature mapping constructed in the Random Fourier Feature method.
\mathcal{D}	The dataset, $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$.
$\text{diag}(\mathbf{w})$	Vector argument, a diagonal matrix containing the elements of vector \mathbf{w} .
$\text{diag}(\mathbf{W})$	Matrix argument, a vector containing the diagonal elements of the matrix \mathbf{W} .
\mathbb{E} or $\mathbb{E}_{q(x)}[z(x)]$	Expectation, or expectation of $z(x)$ where $x \sim q(x)$.
\mathcal{GP}	Gaussian process $f \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$, the function f is distributed as a Guassian process with mean function $m(\mathbf{x})$ and covariance function $k(\mathbf{x}, \mathbf{x}')$.

$k(\cdot, \cdot)$	A covariance or kernel matrix.
$K_{WW'}$	For two data sets $\mathbf{W} = [\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_n]^\top \in \mathbb{R}^{n \times d}$ and $\mathbf{W}' = [\mathbf{w}'_1, \mathbf{w}'_2, \dots, \mathbf{w}'_m]^\top \in \mathbb{R}^{n' \times d}$ the matrix $K_{WW'} \in \mathbb{R}^{n \times n'}$ has elements $(K_{WW'})_{i,j} = k(\mathbf{w}_i, \mathbf{w}'_j)$.
lin-solve (\mathbf{A}, \mathbf{B})	A function used to solve $\mathbf{X} = \mathbf{A}^{-1}\mathbf{B}$ in the linear system $\mathbf{AX} = \mathbf{B}$.
$\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ or $\mathcal{N}(x \boldsymbol{\mu}, \boldsymbol{\Sigma})$	(the variable x has a) Multivariate Gaussian distribution with mean vector $\boldsymbol{\mu}$ and covariance $\boldsymbol{\Sigma}$.
n and n_*	The number of training (and tests) cases.
N	The dimension of the feature space.
\mathbb{N}	The natural numbers, $\mathbb{N} = \{1, 2, 3, \dots\}$.
$\mathcal{O}(\cdot)$	Big-O notation. If a function $f \in \mathcal{O}(g)$ then the absolute value of $f(x)$ is at most a positive multiple of $g(x)$ for all sufficiently large values of x .
$y x$ and $p(x y)$	A conditional random variable y given x and its probability density.
\mathbf{Q}, \mathbf{V}	Typically used to denote a matrix with orthonormal structure.
\mathbb{R}	The real numbers.
$\text{tr}(\mathbf{A})$	The trace of a matrix.
\mathbb{V} or $\mathbb{V}_{q(x)}[z(x)]$	Variance, the variance of $z(x)$ when $x \sim q(x)$.
\mathcal{X}	Input space.
\mathbf{X}	The $n \times d$ matrix of training inputs.
\mathbf{X}_*	The $n_* \times d$ matrix of test inputs.
\mathbf{x}_i	The i^{th} training input.
\mathbb{Z}	The integers, $\mathbb{Z} = \{\dots, -2, -1, 0, 1, 2, \dots\}$.

INTRODUCTION

Time series prediction (and related regressional tasks) is a subject of high interest across many disciplines of science and mathematics. The history of time series can be traced back to the birth of science in ancient Greece where Aristotle devised a systematic approach to weather forecasting in 350 BC in his famous treatise *Meteorologica*. This method was later used to help predict when certain meteorological induced events, such as the flooding of the Nile river [HHF73]. Statistical modelling for time series prediction would not come until the 20th century where development of AutoRegressive Moving Average (ARMA) models which were first mentioned by Yule [Yul27] in 1927 and later popularized by Box and Jenkins in their book *Time Series Analysis* published in 1970 [Box08].

Given a data set of n observations $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$, where each input $x_i \in \mathbb{R}_{>0}$ is a time value and $y_i \in \mathbb{R}$ is a output or experimental observation that acts a function of time, the goal of time series prediction is to try and best predict a value y_* at a novel time x_* . With computing power becoming ever more advanced and affordable, many have turned to Machine Learning (ML) to develop sophisticated models to address the problem of creating accurate yet computationally inexpensive time series predictors. Broadly speaking, ML is any class of heuristic algorithm that attempts to refine and develop some model to perform a "simple" task by learning through user provided input. ML is founded on the idea that any form of task learning is done through sensory input taken from the surrounding environment. More formally speaking, ML attempts to generate a function $f : X \rightarrow Y$, for some input set X and observation or output set Y , where the outputs given by f closely aligns to actual observations. It is tacitly assumed that the phenomena we are studying follows laws which admit mathematical formulation and that experimental results can be reproduced to some degree of accuracy. Typically, experiments will never produce exact values of the underpinning law, g . Instead experimental observations, y_i , will include a small amount of random error so that $y_i = g(x_i) + \varepsilon_i$ where $\varepsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$.

A ML model will attempt to make accurate predictions using some simplified formulation of the world. The distribution corresponding to the probability of a prediction within the context of the "state of the world" is referred to as the *likelihood*. The uncertainty within the likelihood stems from the predictive limits of the model. These limitations usually arise as a consequence of selecting a model which is either too simple or complex. The "state of the world" is sometimes internally captured by the model as a set of mutable parameters θ . The process of taking observations and using them to form predictions is called *inference* which, in some sense, is synonymous with learning [VdW19].

ML can be applied to time series prediction in a fairly straight forward manner by simply teaching a ML algorithm the time series data set, \mathcal{D} , to hopefully produce a function f that serves as a good approximant for event prediction.

In this thesis we shall focus on a particular class of ML algorithms called Bayesian models which, unsurprisingly, makes use of Bayesian statistics to drive inference. In Bayesian models a *prior* distribution is used to quantify the uncertainty of the current state of the model before any observations are made. The model can then be updated once data is observed by using the likelihood to give a *posterior* distribution which represents the reduced uncertainty after "teaching" the model new observations. Methods of teaching a model how to change its behavior using a new set of observations often involves the use of a

loss function L . The loss function is used as an aid in deciding what action, a , should be taken in to best minimize uncertainty. The best action, roughly speaking, can be evaluated as

$$a_{\text{opt}} = \arg \min_a \int L(y_*, a) p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) dy_*.$$

Interestingly, the best action does not rely so much on the model's internalized parameters but rather on the predictive distribution $p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y})$ [VdW19]. This key insight has spawned a class of ML algorithms that focuses on inferring the function f directly by computing $p(f | \mathcal{D})$ instead of finding optimal internal parameters using $p(\theta | \mathcal{D})$ [Mur12]. Models that perform inference in this manner are called *non-parameteric* models. Within the *non-parameteric* model paradigm, the predictive distribution can be represented as

$$p(y_* | x_*, \mathbf{X}, \mathbf{y}) = \int p(y_* | f, x_*) p(f | \mathbf{X}, \mathbf{y}) df$$

and once new data is observed the posterior can be updated using Baye's rule

$$\text{posterior} = \frac{\text{likelihood} \times \text{prior}}{\text{marginal likelihood}}, \quad p(f, f_* | \mathbf{y}) = \frac{p(\mathbf{y} | f) p(f, f_*)}{p(\mathbf{y})}$$

[Ras06]. This thesis will focus on a particular non-parameteric Bayesian ML model called Gaussian processes (GPs). The over arching idea of GPs is to assign a prior probability to every possible function mapping from X to Y . While this does not appear to be computationally tractable due to the seemingly uncountable infinite number of mappings that would require checking, it turns out, these computations can infact be carried out given we are only seeking predictions at a finite number of points using a finite number of observations. GPs occupy a special place within the realm of ML since they account for uncertainty in a principled way, are relatively simple to implement and are highly modular allowing them to easily be incorporated into a larger systems. It is no surprise then that while other kernel methods (such as kernelized k^{th} nearest neighbors and ridge regression) are still overshadowed by their neural network cousins, GPs have made a quiet comeback in the ML community [Cao18].

The following example highlights a particular GP success story: a team of researchers led by Andries Potgieter at QAAFI (UQ) are currently investigating new digital approaches to accurately derive crop phenology stages (i.e. mid green, peak, flowering, grain filling and harvest) measured at field scale across large regions. Such methods can be used to better inform farmers and industry on the optimised time to plant various crops to minimize crop loss from environmental stresses such as frost and fungal disease. This involves analysing crop growth from previous seasons (i.e. 2018-2021) to forecast when certain phenological stages will take place in the current harvest. Outputs form this tool will allow producers to accurately map the temporal and spatial extend of phenology at a field and farm levels across different regions and seasons. This problem is readily converted into a time series problem. Originally, Potgieter's team surveyed a number of different parameteric models to carry out forecasting. However, the parameteric models we serverely limited in their ability to inform when key phenological stages would take place. After seeing the success of applying GPs to other remote sensing tasks [SD22] the team investigated the use of GPs in their own research to find that they could produce much higher resolution predictions from which they could infer a far richer phenological timeline [Pot13]. A comparision of using GPs over other parameteric models is shown in Figure 1. Potgieter's team found that the only draw back to using GPs was the lengthy run time required to create predictions and fears that

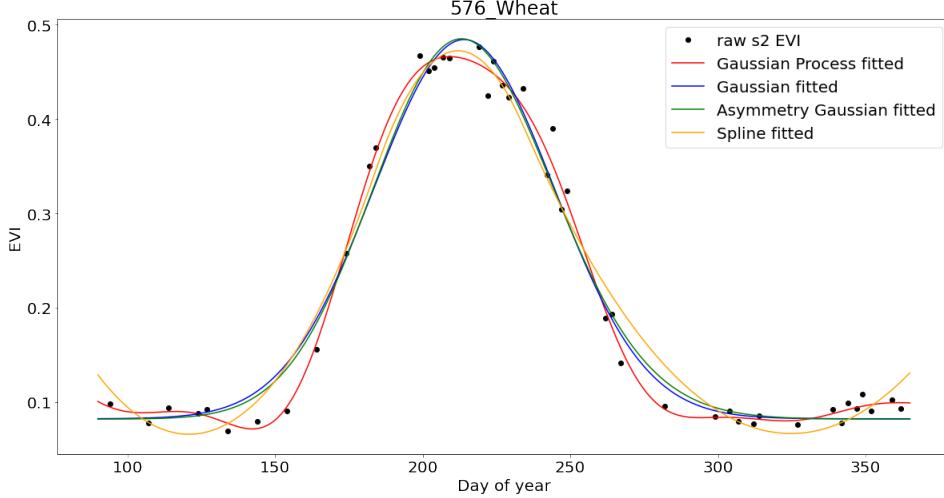


FIGURE 1. Potgieter’s team found that GPs were superior in terms of predicting a phenological timeline for a number of common seasonal crops over other parametric models.

collecting new data each season will only exacerbate the issue. This is a common problem shared by anyone wanting to use GPs. Due to their unwieldy $\mathcal{O}(n^3)$ runtime, where n is the number of observations, GPs become impractical to apply on datasets with $n > 10^5$ samples. As such, the goal of this thesis is to explore various avenues one can take to replace some of the more intense calculations of GPs with computationally more efficient approximations without overly sacrificing accuracy.

Chapter ?? will give a more mathematical treatment of GPs starting from the ground through a review of some fundamental material from functional analysis also the theory behind the motivation of GPs before finally concluding with concrete algorithms for GP regression and classification. Chapters ?? and ?? will cover techniques for approximating a large matrix used with GPs that provides information on how similar each observation is to one other. Chapter ?? then gives alternative methods for solving linear systems, an essential component required for the GP algorithm to work.

GAUSSIAN PROCESSES

The aim of this chapter is to explore the theory behind GPs. First, some essential theory from functional analysis on kernels and reproducing kernel Hilbert spaces will be reviewed which are not only used in GPs but are found in a vast array of machine learning models, aptly named kernel machines. Afterward, we shall go through the underlying statistics that drive GP prediction and use it to form algorithms for both regression and classification tasks. Note that most of the theory presented here is only for real-values data sets although most the time complex-valued generalizations do exist.

0.1. Kernels. Often in machine learning we are often met with the challenge of how to best represent data instances as fixed size feature vectors $\mathbf{x}_i \in X$. For certain objects it might not be obvious at all how to represent the data as a fixed length vector. Good examples of variable length data include textual documents and genomic data. For these data types we can define a method of measuring similarity between objects which requires them to first be converted to a fixed length feature vector first [Mur12]. To do this we begin by mapping the feature vectors into a Hilbert space H which enriches the vector space with an inner product $\langle \cdot, \cdot \rangle_H : H \times H \rightarrow \mathbb{R}$ and a norm $\|\cdot\|_H : H \rightarrow \mathbb{R}$. Input data is transformed into feature space vectors via a non-linear feature mapping $\Phi : X \rightarrow H$. The benefit of using feature maps in this way is that a non-linear descision boundary can be constructed using linear models. In some instances a similarity measure can be computed directly using a function $k : X \times X \rightarrow \mathbb{R}$, instead of needing to construct a Φ and then computing the inner product of the transformed instances. Functions that act directly on our data instances are known as kernel functions and using them to avoid computation associated with the underlying feature space is known as the kernel trick [Ste08]. These ideas are stated more formally in definition 1.

Definition 1 (Kernel). Let X be a non-empty set. Then a function $k : X \times X \rightarrow \mathbb{R}$ is called a kernel on X if there exists a Hilbert space and a map $\Phi : X \rightarrow H$ such that for all $\mathbf{x}, \mathbf{x}' \in X$ we have $k(\mathbf{x}, \mathbf{x}') = \langle \Phi(\mathbf{x}), \Phi(\mathbf{x}') \rangle_H$. We call the Φ the feature map and H the feature space of k .

It is worth noting that almost no conditions are placed on the set X , allowing it to accommodate virtually any form of data. It is not surprising then that neither the feature map nor the feature space are uniquely determined by the kernel. As shown by the example from Steinwart and Christmann [Ste08], when $X = \mathbb{R}$ and $k(x, x') = x \cdot x'$ where $x, x' \in X$, we can see that k is a kernel using the feature map $\Phi(x) = x$ and $H = \mathbb{R}$. However, another suitable feature map for this particular kernel is $\Phi'(x) = (x/\sqrt{2}, x/\sqrt{2})$ with a corresponding feature space of $H = \mathbb{R}^2$ since

$$\langle \Phi'(x), \Phi'(x') \rangle_{\mathbb{R}^2} = \frac{x'}{\sqrt{2}} \cdot \frac{x}{\sqrt{2}} + \frac{x'}{\sqrt{2}} \cdot \frac{x}{\sqrt{2}} = x \cdot x'$$

for $x, x' \in X$. While their might be numerous functions that provide some notion of similarity between data entries, these functions might not be valid kernels. Instead of needing to construct a feature map and feature space to verify that a chosen function is a valid kernel using definition 1, we can make use of a much simpler set of criteria. Before embarking on this train of thought, we need to define the following.

Definition 2 (Positive Definite and Positive Semidefinite). A function $k : K \times K \rightarrow \mathbb{R}$ is positive semidefinite if for all $n \in \mathbb{N}$ and $\alpha_1, \dots, \alpha_n \in \mathbb{R}$ and all $\mathbf{x}_1, \dots, \mathbf{x}_n \in X$ we have

$$(1) \quad \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j k(\mathbf{x}_j, \mathbf{x}_i) \geq 0.$$

Furthermore, k is said to be positive definite if for mutually distinct $\mathbf{x}_1, \dots, \mathbf{x}_n \in X$ equality 1 only holds for $\alpha_1 = \dots = \alpha_n = 0$ [Ste08].

Definition 3 (Symmetric). A function $k : K \times K \rightarrow \mathbb{R}$ is called symmetric if $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$ for any inputs $\mathbf{x}', \mathbf{x} \in X$ [Ste08].

Definition 4 (Gram Matrix). For fixed $\mathbf{x}_1, \dots, \mathbf{x}_n \in X$ the matrix $\mathbf{K} \in \mathbb{R}^{n \times n}$ where $\mathbf{K}_{i,j} \triangleq k(\mathbf{x}_j, \mathbf{x}_i)$ is the Gram matrix [Ste08].

Note that checking if a function is positive (semi) definite is equivalent to checking that any Gram matrix produced by a function is positive (semi) definite. If k is a real valued kernel corresponding to the feature map Φ , then k is symmetric by virtue of the fact that the inner product of a real Hilbert space is symmetric. Moreover k is positive definite since for $\alpha_1, \dots, \alpha_n \in \mathbb{R}$ and $\mathbf{x}_1, \dots, \mathbf{x}_n \in X$ we have

$$\begin{aligned} & \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j k(\mathbf{x}_j, \mathbf{x}_i) \\ &= \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j \langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle_H \\ &= \left\| \sum_i^n \alpha_i \Phi(\mathbf{x}_i) \right\|_H^2 \\ &\geq 0. \end{aligned}$$

The following theorems tell us that it is not only necessary for a kernel to be positive semi definite but it is also a sufficient condition.

Theorem 5. A function $k : K \times K \rightarrow \mathbb{R}$ is a kernel if and only if it is symmetric and positive semidefinite [Ste08].

0.2. Reproducing Kernel Hilbert Spaces. We shall now shift our attention towards reproducing kernel Hilbert spaces (RKHS) and describe their relation to kernels, and see that in some sense, the RKHS of a kernel k is the smallest feature space for a kernel. The formal definition of a RKHS is stated in definition 6.

Definition 6 (RKHS). Let $X \neq \emptyset$ and H be a real Hilbert space over X

- (1) A function $k : X \times X \rightarrow \mathbb{R}$ is called a reproducing kernel if we have $k(\cdot, \mathbf{x}) \in H$ for all $\mathbf{x} \in X$ and the reproducing property

$$f(\mathbf{x}) = \langle f, k(\cdot, \mathbf{x}) \rangle$$

holds for all $f \in H$ and $\mathbf{x} \in X$.

- (2) The space H is called a reproducing kernel Hilbert space over X if for all $\mathbf{x} \in X$ the Dirac functional $\delta_{\mathbf{x}} : H \rightarrow \mathbb{R}$ defined by $\delta_{\mathbf{x}}(f) \triangleq f(\mathbf{x})$, $f \in H$ is continuous.

[Ste08]

An important property of the RKHS is that the convergence in the norm implies pointwise convergence. Specifically, in a RKHS for any sequence of functions $\{f_n\} \subset H$ where $\|f_n - f\| \rightarrow 0$ we have $|\delta_{\mathbf{x}}(f_n) - \delta_{\mathbf{x}}(f)| = |f_n(\mathbf{x}) - f(\mathbf{x})| \rightarrow 0$. Note that because the evaluation function is both linear and continuous, then it is also bounded in the sense that there is a $c \in \mathbb{R}$, $c > 0$ such that for every $f \in H$ and a fixed $\mathbf{x} \in X$ we have $|\delta_{\mathbf{x}}(f)| \leq c \|f\|_H$ [Ber96]. This property of uniform convergence implying pointwise convergence is important since it tells us that if functions $f, g \in H$ are close in norm then their evaluation at any point is also similar. The following lemma ties together the definition of an RKHS, reproducing kernel and a kernel.

Lemma 7. *Let H be a Hilbert function space over X that has a reproducing kernel k . Then H is a RKHS and H is also a feature space of k where the feature map $\Phi : X \rightarrow H$ is given by*

$$\Phi(\mathbf{x}) = k(\cdot, \mathbf{x})$$

for some $\mathbf{x} \in X$. We call Φ the canonical feature map.

Proof. Since the reproducing property tells us that any Dirac functional can be represented by the reproducing kernel this means

$$|\delta_{\mathbf{x}}(f)| = |f(\mathbf{x})| = |\langle f, k(\cdot, \mathbf{x}) \rangle| \leq \|k(\cdot, \mathbf{x})\|_H \cdot \|f\|_H$$

for all $\mathbf{x} \in X$, $f \in H$. This shows continuity of $\delta_{\mathbf{x}}$ for $\mathbf{x} \in X$. In order to show that Φ is a feature map, fix an $\mathbf{x}' \in X$ and set $f = k(\cdot, \mathbf{x}')$. Then for $\mathbf{x} \in X$, the reproducing property yields

$$\langle \Phi(\mathbf{x}'), \Phi(\mathbf{x}) \rangle_H = \langle k(\cdot, \mathbf{x}'), k(\cdot, \mathbf{x}) \rangle_H = \langle f, k(\cdot, \mathbf{x}) \rangle_H = f(\mathbf{x}) = k(\mathbf{x}', \mathbf{x}).$$

□

This tells us that every Hilbert space with a reproducing kernel is a RKHS. We can also show the converse, that is, every RKHS has a unique reproducing kernel seen in theorem 8.

Theorem 8. *Let H be a RKHS over X . Then $k : X \times X \rightarrow \mathbb{R}$ defined by $k(\mathbf{x}', \mathbf{x}) = \langle \delta_{\mathbf{x}}, \delta_{\mathbf{x}'} \rangle_H$, $\mathbf{x}, \mathbf{x}' \in X$ is the only reproducing kernel of H [Ste08].*

Theorem 8 shows that a RKHS is uniquely determined by its kernel. In fact the other direction can also be shown to afford a one-to-one correspondence between kernels and RKHS. This is known as the Moore-Aronszajn theorem presented in theorem 9.

Theorem 9 (Moore-Aronszajn). *Suppose k is a symmetric positive definite kernel on a set X . Then there is a unique Hilbert space of functions for which k is the reproducing kernel [Ber03].*

The elements of a RKHS will often inherit the analytical properties of its corresponding kernel. This means that kernels provide a mechanism for generating spaces of functions with useful analytical properties.

0.3. Gaussian Radial Basis Kernel. Although there are many kernels to use at our disposal, we turn our attention to a specific class of kernel that shall be used extensively in the upcoming theory and experimentation.

Definition 10 (Gaussian Radial Basis Kernel). *For $d \in \mathbb{N}$, $\sigma \in \mathbb{R}_{>0}$ and $\mathbf{z}, \mathbf{z}' \in \mathbb{R}^d$ we define*

$$k_\sigma(\mathbf{z}, \mathbf{z}') \triangleq \exp\left(-\sigma^{-2} \sum_{j=1}^d (z_j - z'_j)^2\right).$$

Then k_σ is a real valued kernel called the Gaussian Radial Basis Kernel (RBF) kernel with bandwidth σ . Moreover k_σ can be computed as

$$\exp\left(\frac{-\|\mathbf{z} - \mathbf{z}'\|_2^2}{\sigma^2}\right)$$

[Ste08].

The Gaussian RBF kernel makes for a very simple and intuitive measurement of similarity between its inputs. One geometric interpretation of the Gaussian RBF kernel is that as the radius of the smallest d -sphere containing $\mathbf{z}, \mathbf{z}' \in \mathbb{R}^d$ grows the corresponding measurement of similarity decays exponentially. A visual representation of this decay is shown in Figure 2.

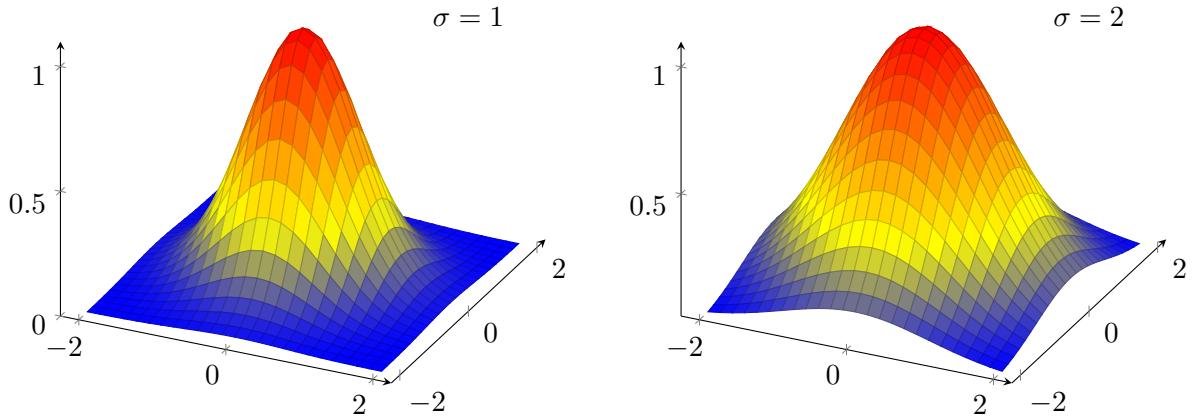


FIGURE 2. A graph of the Gaussian RBF from definition 10 for $d = 2$. Evidently, a larger value of σ slows the rate of decay increasing the similarity between the same pair of samples.

This kernel is infinitely differentiable meaning it has mean square derivatives of all orders and is therefore very smooth. In fact, some argue that such strong smoothness makes it unrealistic for modelling natural phenomena [Ras06, Ste99]. Nonetheless, Gaussian RBF kernels remain the one of the most widely used in literature.

0.4. Kernel Machines.

0.4.1. *Introduction to Support Vector Machines for Binary Classification.* In this section, we will be investigating at two different machine learning models that make use of kernels to perform classification and regression. The first class of kernel machines to be evaluated are support vector machines (SVM). SVMs were originally designed for binary classification and as such only a model for binary classification is presented, although extensions exist that allow regression and multi-class classification.

For the binary classification problem we are tasked with labelling new samples with either one of two classes, -1 or 1 . We shall assume our input space consists of vectors from \mathbb{R}^d and that we provided with a labelled training set $D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$. One simple method to classify samples is by creating an affine linear hyperplane satisfying

$$(2) \quad \begin{aligned} \langle \mathbf{w}, \mathbf{x}_i \rangle + b &> 0, & y_i &= +1 \\ \langle \mathbf{w}, \mathbf{x}_i \rangle + b &< 0, & y_i &= -1 \end{aligned}$$

for some $\mathbf{w} \in \mathbb{R}^d$ and $b \in \mathbb{R}$ where $\|\mathbf{w}\|_2 = 1$ [Ste08]. Moreover we would like \mathbf{w} and b to maximise the margin, that is the maximal distance between the separating hyperplane and the points in D . The specific \mathbf{w} and b obtained through the training set is denoted \mathbf{w}_D and b_D and the resulting decision function is defined as

$$f_D(\mathbf{x}) \triangleq \text{sign}(\langle \mathbf{w}_D, \mathbf{x} \rangle + b_D).$$

There are, however, a number of shortcomings to this model. The most obvious is that our training data may not be linearly separable in \mathbb{R}^d meaning no such \mathbf{w}_D and b_D exist. Moreover, when noise is introduced to the data set this model will prioritize finding a hyperplane that perfectly separates the two classes, making no compromises in misclassifying points, and consequently subjecting it to overfitting. SVMs were introduced by Boser *et al.* [Bos92] to address the first issue of separability. Their approach was to lift the input vector into a more malleable Hilbert space H_0 using a feature map. The inputs are then classified within the new space. Unfortunately this method does nothing to address the second issue of over fitting and, if anything, actually worsens it. Cortes and Vapnik [Cor95] attempted to address this second issue by introducing slack variables to equation 2 so that now we need to satisfy $y_i (\langle \mathbf{w}, \Phi(\mathbf{x}_i) \rangle + b) \geq 1 - \xi_i$ for some $\xi_i \in \mathbb{R}_{\geq 0}$. These constraints can be re-written as

$$\xi_i \geq 1 - y_i (\langle \mathbf{w}, \Phi(\mathbf{x}_i) \rangle + b)$$

and combining this with our slack constraints (that is $\xi_i \geq 0$) yields

$$\xi_i \geq \max\{0, 1 - y_i (\langle \mathbf{w}, \Phi(\mathbf{x}_i) \rangle + b)\} = L_{\text{hinge}}(y_i, \langle \mathbf{w}, \Phi(\mathbf{x}_i) \rangle + b)$$

where L_{hinge} is the hinge loss defined as

$$L_{\text{hinge}}(y, \eta) \triangleq \max\{0, 1 - y\eta\}.$$

This optimization problem can be re-written in the form

$$\min_{(\mathbf{w}, b) \in H_0 \times \mathbb{R}} \lambda \|\mathbf{w}\|_{H_0} + \frac{1}{n} \sum_{i=1}^n L_{\text{hinge}}(y_i, f_{(\mathbf{w}, b)})$$

where $f_{(\mathbf{w}, b)} : X \rightarrow \mathbb{R}$ is defined as

$$f_{(\mathbf{w}, b)} \triangleq \langle \mathbf{w}, \Phi(\mathbf{x}_i) \rangle + b$$

[Ste08]. Unfortunately, this new embedding requires us to solve for optimal parameters in a very high, or even infinite, dimension vector space. To get around this, the Lagrange approach is typically used to solve the corresponding dual problem. When the hinge loss is used the dual problem becomes

$$(3) \quad \begin{aligned} & \max_{\alpha \in [0, C]^n} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n y_i y_j \alpha_i \alpha_j \langle \Phi(\mathbf{x}_i), \Phi(\mathbf{x}_j) \rangle \\ & \text{subject to } \sum_{i=1}^n y_i \alpha_i = 0 \end{aligned}$$

Notice that in the dual problem, we find that inner products are only taken with vectors that have the feature map applied to them allowing us to employ the kernel trick described in section 0.1 so that equation 3 becomes

$$\begin{aligned} & \max_{\alpha \in [0, C]^n} \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i,j=1}^n y_i y_j \alpha_i \alpha_j k(\mathbf{x}_i, \mathbf{x}_j) \\ & \text{subject to } \sum_{i=1}^n y_i \alpha_i = 0. \end{aligned}$$

0.4.2. *Introduction to Gaussian Processes for Regression.* The next machine learning model of interest that uses kernels are gaussian processes. To motivate this model, consider the time series data in Figure 3 (A).

In this diagram there is a number of observed values $D = \{(x_1, y_1), (x_2, y_2), (x_3, y_3), (x_4, y_4)\}$ (blue labels) as well as a missing observation at time x_* . This is a classic problem of time series prediction, that is what is a good prediction for the missing value at time x_* ? Perhaps something close to the red diamond seen in Figure 3 (B)? Why does this red diamond seem like a good choice? Because for known data for which the measurement of similarity is small, we expect the corresponding outputs should also be similar since most natural phenomena are continuous by nature. This reasoning is used to underpin GPs, that is, training inputs that are more similar to the input value should have greater influence over the prediction.

Like SVMs, we can motivate the mathematical model of a GP through a linear model. Since GPs are designed for regression tasks, we shall only focus on GP regression although we will see later that GPs can be extended to perform binary classification and even multi-class classification. To begin, consider the following linear regression model

$$(4) \quad f(\mathbf{x}) \triangleq \langle \mathbf{w}, \mathbf{x} \rangle$$

where we again assuming that \mathbf{x} belongs to \mathbb{R}^d and that $\mathbf{w} \in \mathbb{R}^d$ is a weight vector. Notice the striking resemblances to the linear classifier used in our derivation for the SVM model, although this time we are using the value computed by the inner product directly to infer instead of fitting it over a sign function to force it into a binary class. Suppose we have independently sampled observations $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ to a noisy version of f

$$y_i = f(\mathbf{x}_i) + \varepsilon_i$$

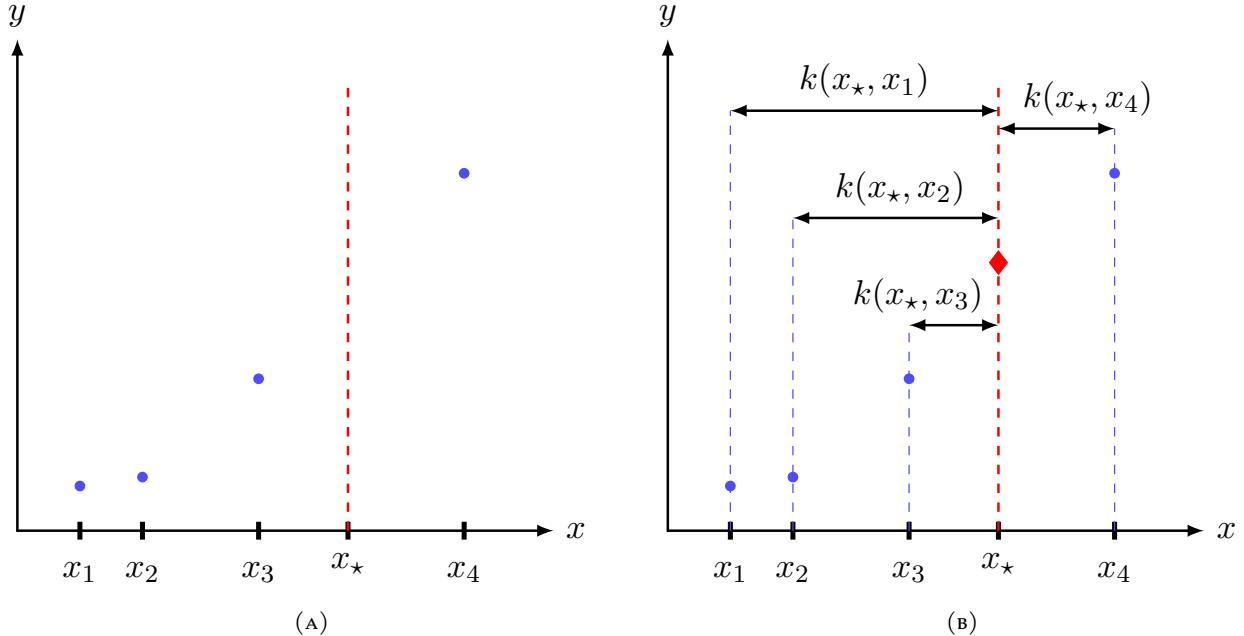


FIGURE 3. Panel (A) shows depicts the classical problem of time series prediction, guessing a value for x_* given values for surrounding times. Panel (B) shows a suitable choice for the value at time x_* with the reasoning that closely surrounding values should have greater influence over inference. Example taken from the 2013 Machine Learning course CPSC 540 instructed by Nando de Freitas [NdF13].

where $\varepsilon_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma_n^2)$. Together the assumption of noise and the base linear model give rise to a likelihood, or more specifically, a probability density over the observations given the inputs and weight parameters. Due to the assumption of independence in our observations

$$\begin{aligned}
(5) \quad p(y \mid \mathbf{X}, \mathbf{w}) &= \prod_{i=1}^n p(y_i \mid \mathbf{x}_i, \mathbf{w}) \\
&= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma_n} \exp\left(-\frac{(y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle)^2}{2\sigma_n^2}\right) \\
&= \frac{1}{(2\pi\sigma_n^2)^{\frac{n}{2}}} \exp\left(-\frac{1}{2\sigma_n^2} \left(\sum_{i=1}^n (y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle)^2\right)\right) \\
&= \frac{1}{(2\pi\sigma_n^2)^{\frac{n}{2}}} \exp\left(-\frac{1}{2\sigma_n^2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2\right) \\
&= \mathcal{N}(\mathbf{X}\mathbf{w}, \sigma_n^2 \mathbb{1}_{n \times n})
\end{aligned}$$

where $\mathbf{y} = [y_1, y_2, \dots, y_n]^\top \in \mathbb{R}^n$ and $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]^\top \in \mathbb{R}^{n \times d}$ [Ras06, page 9]. Within the Bayesian paradigm, a prior is required to represent our beliefs about the parameters in the absence of any information. Typically, the following prior is used for the weight vector

$$\boldsymbol{w} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma}_p)$$

where Σ_p is an appropriate covariance matrix. Ideally, we would like to know the posterior pdf $p(\mathbf{w} | \mathbf{y}, \mathbf{X})$ which refines our choices of \mathbf{w} by taking into account our observations. The posterior can be computed using Bayes rule

$$p(\mathbf{w} | \mathbf{y}, \mathbf{X}) \propto p(\mathbf{y} | \mathbf{w}, \mathbf{X}) p(\mathbf{w}).$$

Equation 5 gives us a probability for $p(\mathbf{y} | \mathbf{w}, \mathbf{X})$ and since $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \Sigma_p)$ then

$$p(\mathbf{w}) = \frac{1}{\sqrt{(2\pi)^d |\Sigma|}} \exp\left(-\frac{1}{2}\mathbf{w}^\top \Sigma_p^{-1} \mathbf{w}\right)$$

[Kro14]. This means, up to proportionality

$$\begin{aligned} p(\mathbf{w} | \mathbf{y}, \mathbf{X}) &\propto \exp\left(-\frac{1}{2\sigma_n^2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w})\right) \exp\left(-\frac{1}{2}\mathbf{w}^\top \Sigma_p^{-1} \mathbf{w}\right) \\ &\propto \exp\left(-\frac{1}{2} (\mathbf{w} - \bar{\mathbf{w}})^\top \left(\frac{1}{\sigma_n^2} \mathbf{X}^\top \mathbf{X} + \Sigma_p^{-1}\right) (\mathbf{w} - \bar{\mathbf{w}})\right) \end{aligned}$$

where $\bar{\mathbf{w}} \triangleq \sigma_n^{-2} (\sigma_n^{-2} \mathbf{X}^\top \mathbf{X} + \Sigma_p^{-1})^{-1} \mathbf{X}^\top \mathbf{y}$. Notice that this again is a multivariate Gaussian distribution with mean $\bar{\mathbf{w}}$ and covariance \mathbf{A}^{-1} where $\mathbf{A} \triangleq \sigma_n^{-2} \mathbf{X}^\top \mathbf{X} + \Sigma_p^{-1}$ so that

$$p(\mathbf{w} | \mathbf{X}, \mathbf{y}) \sim \mathcal{N}(\mathbf{w}, \mathbf{A}^{-1}).$$

To make a prediction of our target function for an input, \mathbf{x}_* , outside our observed values we can take the average over all possible parameter values weighted by the posterior to predict $f_* \triangleq f(\mathbf{x}_*)$ which yields

$$p(f_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int_{\mathbb{R}^d} p(f_* | \mathbf{x}_*, \mathbf{w}) p(\mathbf{w} | \mathbf{X}, \mathbf{y}) d\mathbf{w} = \mathcal{N}(\mathbf{x}_*^\top \bar{\mathbf{w}}, \mathbf{x}_*^\top \mathbf{A}^{-1} \mathbf{x}_*).$$

This gives another Gaussian distribution whose mean is the mean of the posterior distribution of the weight vectors multiplied by the input vector, and whose covariance in the quadratic form of the covariance of the weight vectors again with the input vectors. This makes sense since it tells us that the uncertainty of the model grows quadratically with the magnitude of the input.

We can now employ the kernel trick in the exact same manner in the derivation of the SVM model, that is, by using a feature mapping Φ to lift the inputs of our linear regression model from equation 4 into a higher dimension and more workable Hilbert space so that our model now becomes

$$f(\mathbf{x}) \triangleq \langle \mathbf{w}, \Phi(\mathbf{x}) \rangle.$$

The derivation for the new model is identical with the only difference being that \mathbf{x}_* is replaced with $\Phi(\mathbf{x}_*)$ and \mathbf{X} is replaced with $\Phi(\mathbf{X}) \triangleq [\Phi(\mathbf{x}_1), \Phi(\mathbf{x}_2), \dots, \Phi(\mathbf{x}_n)]^\top \in \mathbb{R}^{n \times N}$ where N is the dimension of the feature space. The new predictive distribution can be expressed as

$$(6) \quad f_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y} \sim \mathcal{N}\left(\frac{1}{\sigma_n^2} \Phi(\mathbf{x}_*)^\top \mathbf{A}^{-1} \Phi(\mathbf{X})^\top \mathbf{y}, \Phi(\mathbf{x}_*)^\top \mathbf{A}^{-1} \Phi(\mathbf{x}_*)\right)$$

where \mathbf{A} is now $\mathbf{A} \triangleq \frac{1}{\sigma_n^2} \Phi(\mathbf{X})^\top \Phi(\mathbf{X}) + \Sigma_p^{-1} \in \mathbb{R}^{N \times N}$. From this, it becomes evident that the inverse of \mathbf{A} is required to compute both the mean and the covariance. This is not favourable since this would require knowledge of the feature space into which the feature map sends inputs. Moreover, computing \mathbf{A}^{-1} may become impractical in the dimension of the feature space is incredibly large. Remember, the whole point of the kernel trick is to avoid any computation that involves direct knowledge of H but rather to use a kernel k to bypass these obstacles and indirectly produce inner products of the data applied to the

feature map. With this in mind, let us try and find different expressions for the mean and the covariance of equation 6 that will enable us to apply the kernel trick. Before starting, we need to find a suitable expression for the mean. First define the notation

$$\mathbf{K}_{WW'} \triangleq \Phi(\mathbf{W}) \Sigma_p \Phi(\mathbf{W}')^\top \in \mathbb{R}^{n \times n'}$$

where $\mathbf{W} \in \mathbb{R}^{n \times d}$ and $\mathbf{W}' \in \mathbb{R}^{n' \times d}$ are two data matrices. Consider the following

$$\begin{aligned} & \mathbf{A} \Sigma_p \Phi(\mathbf{X})^\top \\ &= (\sigma_n^{-2} \Phi(\mathbf{X})^\top \Phi(\mathbf{X}) + \Sigma_p^{-1}) \Sigma_p \Phi(\mathbf{X})^\top \\ &= \sigma_n^{-2} \Phi(\mathbf{X})^\top \Phi(\mathbf{X}) \Sigma_p \Phi(\mathbf{X})^\top + \Phi(\mathbf{X})^\top \\ &= \sigma_n^{-2} \Phi(\mathbf{X})^\top (\Phi(\mathbf{X}) \Sigma_p \Phi(\mathbf{X})^\top + \sigma_n^2 \mathbb{1}_{n \times n}) \\ &= \sigma_n^{-2} \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n}) \end{aligned}$$

meaning

$$\begin{aligned} \sigma_n^{-2} \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n}) &= \mathbf{A} \Sigma_p \Phi(\mathbf{X})^\top \\ \sigma_n^{-2} \mathbf{A}^{-1} \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n}) &= \Sigma_p \Phi(\mathbf{X})^\top \\ \sigma_n^{-2} \mathbf{A}^{-1} \Phi(\mathbf{X})^\top &= \Sigma_p \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n})^{-1} \end{aligned}$$

so that the current mean of

$$\frac{1}{\sigma_n^2} \Phi(\mathbf{x}_*)^\top \mathbf{A}^{-1} \Phi(\mathbf{X})^\top \mathbf{y}$$

in equation 6 can be replaced with

$$\Phi(\mathbf{x}_*)^\top \Sigma_p \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n})^{-1} \mathbf{y}.$$

To find a more suitable expression for the covariance matrix, we will need the assistance of the matrix inversion lemma stated without proof in lemma 11.

Lemma 11 (Matrix Inversion Lemma). *For $\mathbf{Z} \in \mathbb{K}^{n \times m}$, $\mathbf{W} \in \mathbb{K}^{m \times m}$ and $\mathbf{U}, \mathbf{V} \in \mathbb{K}^{n \times m}$ then*

$$(\mathbf{Z} + \mathbf{U} \mathbf{W} \mathbf{V}^\top)^{-1} = \mathbf{Z}^{-1} - \mathbf{Z}^{-1} \mathbf{U} (\mathbf{W}^{-1} + \mathbf{V}^\top \mathbf{Z}^{-1} \mathbf{U})^{-1} \mathbf{V}^\top \mathbf{Z}^{-1}$$

assuming the relevant inverses exist [Pre92, page 75].

Consider

$$(7) \quad \mathbf{A} = \Sigma_p^{-1} + \Phi(\mathbf{X})^\top (\sigma_n^{-2} \mathbb{1}_{n \times n}) \Phi(\mathbf{X})$$

then applying the matrix inversion lemma by setting $\mathbf{Z}^{-1} = \Sigma_p$, $\mathbf{W}^{-1} = \sigma_n^2 \mathbb{1}_{n \times n}$ and $\mathbf{V} = \mathbf{U} = \Phi(\mathbf{X})$ equation 7 then becomes

$$\Sigma_p - \Sigma_p \Phi(\mathbf{X})^\top (\sigma_n^2 \mathbb{1}_{n \times n} + \Phi(\mathbf{X}) \Sigma_p \Phi(\mathbf{X})^\top)^{-1} \Phi(\mathbf{X}) \Sigma_p.$$

Thus equation 6 can be equivalently formulated as

$$(8) \quad f_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y} \sim \mathcal{N}(\Phi(\mathbf{x}_*)^\top \Sigma_p \Phi(\mathbf{X})^\top (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n})^{-1} \mathbf{y}, \Phi(\mathbf{x}_*)^\top \Sigma_p \Phi(\mathbf{x}_*) - \Phi(\mathbf{x}_*)^\top \Sigma_p \Phi(\mathbf{X})^\top (\sigma_n^2 \mathbb{1}_{n \times n} + \Phi(\mathbf{X}) \Sigma_p \Phi(\mathbf{X})^\top)^{-1} \Phi(\mathbf{X}) \Sigma_p \Phi(\mathbf{x}_*)).$$

The astute reader may have noticed the very suggestive notation of labelling matrices of the form $\Phi(\mathbf{W}) \Sigma_p \Phi(\mathbf{W}')^\top$ as $\mathbf{K}_{\mathbf{W}\mathbf{W}'}$ as though it may have some sort of connection to a kernel. To make this even more obvious, notice that each occurrence of the feature map in both expressions for the mean and covariance in equation 8 can be replaced with a $\mathbf{K}_{\mathbf{W}\mathbf{W}'}$ for some appropriate choice of \mathbf{W} and \mathbf{W}' giving a more notationally cleaner expression

$$(9) \quad f_\star | \mathbf{x}_\star, \mathbf{X}, \mathbf{y} \sim \mathcal{N}(\mathbf{K}_{\mathbf{x}_\star \mathbf{X}} (\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbb{1}_{n \times n})^{-1} \mathbf{y}, k(\mathbf{x}_\star, \mathbf{x}_\star) - \mathbf{K}_{\mathbf{x}_\star \mathbf{X}} (\sigma_n^2 \mathbb{1}_{n \times n} + \mathbf{K}_{\mathbf{XX}})^{-1} \mathbf{K}_{\mathbf{x}_\star \mathbf{X}}^\top).$$

To get a better idea of the connection to kernels, since Σ_p is a symmetric positive semi definite matrix, it defines an inner product

$$\langle \mathbf{x}, \mathbf{y} \rangle_{\Sigma_p} = \mathbf{y}^* \Sigma_p \mathbf{x}, \quad \mathbf{x}, \mathbf{y} \in \mathbb{K}^N$$

[Wan, page 34] so that

$$(10) \quad (\mathbf{K}_{\mathbf{WW}'})_{ij} = \langle \Phi(\mathbf{w}_i), \Phi(\mathbf{w}_j) \rangle_{\Sigma_p} = k(\mathbf{w}_i, \mathbf{w}_j)$$

where k is the kernel with feature map Φ and inner product $\langle \cdot, \cdot \rangle_{\Sigma_p}$. In fact when $\mathbf{W} = \mathbf{W}'$ equation 10 is exactly the Gram matrix with said kernel. Thus GPs are another great example of models that take advantage of the kernel trick. We shall see in the coming chapters on how exactly we can compute predictions at novel points.

0.5. Gaussian Processes for Regression. We saw in section 0.4 that, unlike most other machine learning models, GPs infer over a distribution of functions $p(f | \mathcal{D})$ instead of a vector of parameteric values $p(\boldsymbol{\theta} | \mathcal{D})$. Naively, one may attempt to find a suitable f by fixing a class of functions \mathcal{F} and then search over this class to find a function that best represents the data. However, this may not work well if there is not enough richness in \mathcal{F} to represent the data. Instead, a suitable f is selected by first assigning a prior probability to every possible function using the training data and then to select the function with the highest probability. To keep this computation tractable we only evalute our predicted function at a finite number of points. The prediction itself is found by taking the mean over all functions with respect to the posterior conditioned on the observed data which is assumed to be jointly Gaussian with the input value. This gives rise to Gaussian Process more formally stated in definition 12.

Definition 12 (Gaussian Process). *A Gaussian Process (GP) is a collection of random variables with index set I , such that every finite subset of random variables has a joint Gaussian distribution* [Ras06, Mur12].

A GP is completely characterized by a mean function $m(\mathbf{x})$ and a kernel, which in the context of GPs is sometimes called a covariance function, $k(\mathbf{x}, \mathbf{x}')$ on a real process as

$$\begin{aligned} m(\mathbf{x}) &= \mathbb{E}[f(\mathbf{x})] \\ k(\mathbf{x}, \mathbf{x}') &= \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))]. \end{aligned}$$

GPs define a prior over all possible functions which can be used to create a posterior once enough data has been observed. The prior is used to represent the functions we expect to see before any observations are made. Although defining a prior over all possible functions may seem computationally intractable, we actually only need to define a distribution over a finite number of points. Before any observations

are made, we typically assume that the mean function is the constant zero function, that is $m(\mathbf{x}) = 0$. A function $f(\mathbf{x})$ sampled from a GP with mean $m(\mathbf{x})$ and covariance $k(\mathbf{x}, \mathbf{x}')$ is written as

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

Since a GP is a collection of random variables it must satisfy the consistency requirement, that is, observing some of the values should not change the distribution of any small subset of unobserved values. More specifically if

$$(\mathbf{y}_1, \mathbf{y}_2) \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

then

$$\mathbf{y}_1 \sim \mathcal{N}(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_{1,1})$$

$$\mathbf{y}_2 \sim \mathcal{N}(\boldsymbol{\mu}_2, \boldsymbol{\Sigma}_{2,2})$$

where $\boldsymbol{\Sigma}_{1,1}$ and $\boldsymbol{\Sigma}_{2,2}$ are the relevant sub matrices. Again, we shall us the notation that for set of data $\mathbf{W} = [\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_n]^\top \in \mathbb{R}^{n \times d}$ and $\mathbf{W}' = [\mathbf{w}'_1, \mathbf{w}'_2, \dots, \mathbf{w}'_m]^\top \in \mathbb{R}^{n' \times d}$ we use the notation

$$(K_{\mathbf{W}\mathbf{W}'})_{i,j} \triangleq k(\mathbf{w}_i, \mathbf{w}'_j)$$

where $K_{\mathbf{W}\mathbf{W}'} \in \mathbb{R}^{n \times n'}$. The covariance function completely characterized by its kernel. Unless otherwise stated, the kernel or covariance function used in examples and experimentation is the Gaussian RBF kernel, definition 10.

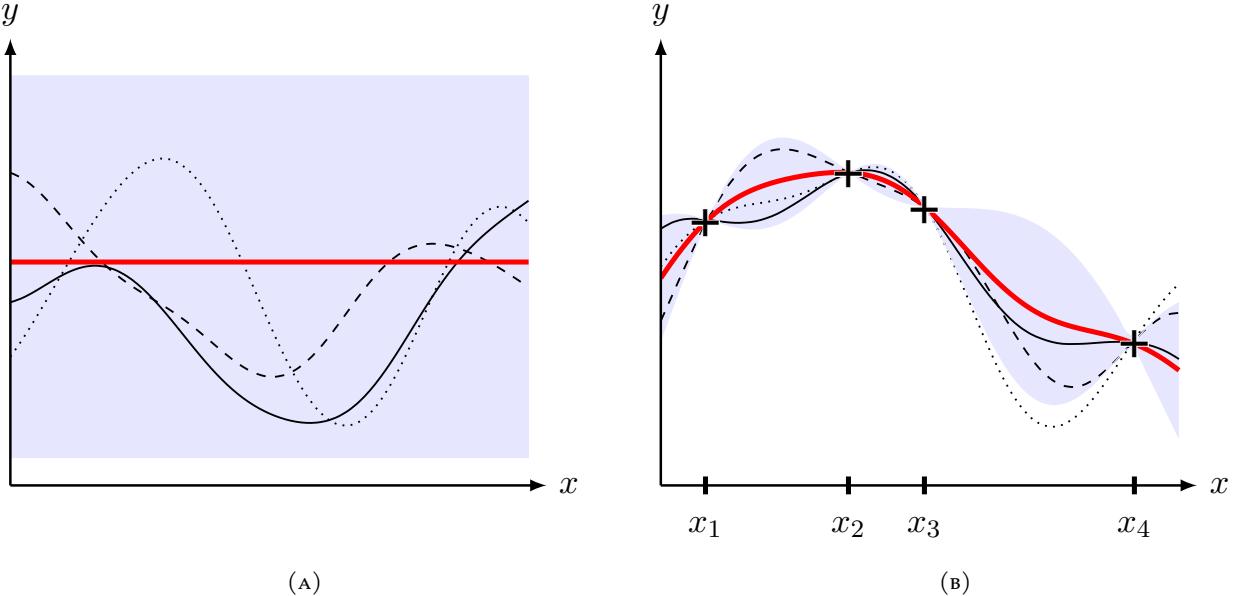


FIGURE 4. Panel (A) shows three function drawn from the prior distribution. Panel (B) shows three function drawn from the prior distribution after four observations have been made. In both panels the mean function is drawn in red, sampled functions in black and twice the standard deviation shaded in light blue.

Figure 4 (A) shows three samples drawn from the prior before any observations are made. GPs also allow us to compute the pointwise variance which can provide some measure of variability for predicted values. The blue shaded area of Figure 4 (A) represents twice the standard deviation about the mean.

0.5.1. *Noise-free observations.* Typically when using GP we would like to incorporate data from observations, or training data, into our predictions on unobserved values. Let us suppose there is some observed data $\mathcal{D} = \{(\mathbf{x}_i, f_i) \mid i \in \{1, 2, \dots, n\}\}$ which is (unrealistically) noise-free that we would like to model as a GP. In other words, for any sample in our dataset we can be certain that the observed value is the true value of the underlying function we wish to model. Then for the observed data

$$\mathbf{f} \sim \mathcal{N}(\mathbf{0}, \mathbf{K}_{XX}).$$

We would then like to make a prediction for unobserved values say $\mathbf{X}_* = [\mathbf{x}_{1*}, \mathbf{x}_{2*}, \dots, \mathbf{x}_{n*}]$ with value f_* as has a distribution of

$$f_* \sim \mathcal{N}(\mathbf{0}, \mathbf{K}_{X_* X_*}).$$

Here \mathbf{f} and f_* are independent but we would like to give them some sort of correlation. We can do this by having them originate from the same joint distribution. According to the prior, we can write the joint distribution of the training points \mathbf{f} and the test points f_* as

$$\begin{bmatrix} \mathbf{f} \\ f_* \end{bmatrix} \sim \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} \mathbf{K}_{XX} & \mathbf{K}_{X_* X}^\top \\ \mathbf{K}_{X_* X} & \mathbf{K}_{X_* X_*} \end{bmatrix}\right).$$

While the above does give us some information that f_* is related to the observed data and the test inputs, it does not provide any method to evaluate f_* . To do this we shall need the assistance of the following theorem.

Theorem 13. (*Marginals and conditionals of an MVN* [Mur12]) Suppose $\mathbf{x} = (\mathbf{x}_1, \mathbf{x}_2)$ is jointly Gaussian with parameters

$$\boldsymbol{\mu} = \begin{bmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \end{bmatrix}, \quad \boldsymbol{\Sigma} = \begin{bmatrix} \boldsymbol{\Sigma}_{1,1} & \boldsymbol{\Sigma}_{1,2} \\ \boldsymbol{\Sigma}_{2,1} & \boldsymbol{\Sigma}_{2,2} \end{bmatrix}$$

then the posterior conditional is given by

$$\begin{aligned} \mathbf{x}_2 \mid \mathbf{x}_1 &\sim \mathcal{N}(\mathbf{x}_2 \mid \boldsymbol{\mu}_{2|1}, \boldsymbol{\Sigma}_{2|1}) \\ \boldsymbol{\mu}_{2|1} &= \boldsymbol{\mu}_2 + \boldsymbol{\Sigma}_{2,1} \boldsymbol{\Sigma}_{1,1}^{-1} (\mathbf{x}_1 - \boldsymbol{\mu}_1) \\ \boldsymbol{\Sigma}_{2|1} &= \boldsymbol{\Sigma}_{2,2} - \boldsymbol{\Sigma}_{2,1} \boldsymbol{\Sigma}_{1,1}^{-1} \boldsymbol{\Sigma}_{1,2}. \end{aligned}$$

Thus, once data has been observed, finding a mean and covariance for f_* involves a direct application of theorem 13 which gives

$$f_* \mid \mathbf{K}_{X_* X}^\top, \mathbf{K}_{XX}, \mathbf{f} \sim \mathcal{N}(\boldsymbol{\mu}_*, \boldsymbol{\Sigma}_*)$$

where

$$\begin{aligned} \boldsymbol{\mu}_* &= \mathbf{0} + \mathbf{K}_{X_* X} \mathbf{K}_{XX}^{-1} (\mathbf{f} - \mathbf{0}) \\ &= \mathbf{K}_{X_* X} \mathbf{K}_{XX}^{-1} \mathbf{f} \end{aligned}$$

and

$$\Sigma_{\star} = \mathbf{K}_{X_{\star} X_{\star}} - \mathbf{K}_{X_{\star} X} \mathbf{K}_{XX}^{-1} \mathbf{K}_{X X_{\star}}$$

meaning we can write a distribution for f_{\star} as

$$(11) \quad f_{\star} | \mathbf{K}_{X_{\star} X}^{\top}, \mathbf{K}_{XX}, f \sim \mathcal{N}(\mathbf{K}_{X_{\star} X} \mathbf{K}_{XX}^{-1} f, \mathbf{K}_{X_{\star} X_{\star}} - \mathbf{K}_{X_{\star} X} \mathbf{K}_{XX}^{-1} \mathbf{K}_{X X_{\star}}^{\top}).$$

Function values from the unobserved inputs X_{\star} , that is f_{\star} , can be estimated using the joint posterior distribution by evaluating the mean of 11. Figure 4 (B) shows these computations given a data set $\mathcal{D} = \{(x_1, y_1), (x_2, y_2), (x_3, y_3), (x_4, y_4)\}$. Notice that the variance tightens around the observed values since (assuming no noise in our data is present) we now know for certain this is how our target function should behave at x_1, x_2, x_3 and x_4 . Clearly, specifying the properties of the prior is important since it fixes the properties of the functions considered during inference.

0.5.2. Prediction with Noisy observations. When attempting to model our value function we usually do not have access to the value function itself but a noisy version thereof, $y = f(x) + \varepsilon$ where $\varepsilon \sim \mathcal{N}(0, \sigma_n^2)$ meaning the prior on the noisy values becomes

$$\text{cov}(\mathbf{y}) = \mathbf{K}_{XX} + \sigma_n^2 \mathbf{I}.$$

The reason why noise is only added along the diagonal follows from the assumption of independence of noise in our data. We can write out the new distribution of the observed noisy values along the points at which we wish to test the underlying function as

$$\begin{bmatrix} \mathbf{y} \\ f_{\star} \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mathbf{0} \\ \mathbf{f}_{\star} \end{bmatrix}, \begin{bmatrix} \mathbf{K}_{XX} + \sigma_n^2 \mathbf{1}_{n \times n} & \mathbf{K}_{X_{\star} X}^{\top} \\ \mathbf{K}_{X_{\star} X} & \mathbf{K}_{X_{\star} X_{\star}} \end{bmatrix} \right).$$

Using a similar we arrive at a similar condition distribution of $f_{\star} | \mathbf{K}_{X_{\star} X}^{\top}, \mathbf{K}_{XX}, f$ we arrive at one of the most fundamental equations for GP regression tasks

$$(12) \quad f_{\star} | \mathbf{K}_{X_{\star} X}^{\top}, \mathbf{K}_{XX}, y \sim \mathcal{N}(\bar{f}_{\star}, \text{cov}(f_{\star}))$$

where

$$(13) \quad \bar{f}_{\star} \triangleq \mathbf{K}_{X_{\star} X} [\mathbf{K}_{XX} + \sigma_n^2 \mathbf{1}_{n \times n}]^{-1} \mathbf{y}$$

$$(14) \quad \text{cov}(f_{\star}) = \mathbf{K}_{X_{\star} X_{\star}} - \mathbf{K}_{X_{\star} X} [\mathbf{K}_{XX} + \sigma_n^2 \mathbf{1}_{n \times n}]^{-1} \mathbf{K}_{X X_{\star}}^{\top}$$

Remarkably, this gives us the the exact same posterior distribution ascertained from the weight space derivation in equation 9. Notice that the prediction of the mean in equation 13 is a linear combination of the observations, sometimes referred to as a *linear predictor*. Another way of looking at the prediction is seeing it as a linear combination of n kernel evaluations centered at the input x_{\star}

$$\mathbf{f}_{\star} = \sum_{i=1}^n \alpha_i k(\mathbf{x}_i, \mathbf{x}_{\star})$$

where $\boldsymbol{\alpha} = [\mathbf{K}_{XX} + \sigma_n^2 \mathbf{1}_{n \times n}]^{-1} \mathbf{y}$. Intuitively, this expression can be understood by realising that, despite defining the GP using a joint Gaussian distribution over the observations, when making predictions GPs only care about the $(n+1)$ -dimension distribution defined by the n observations and the single test

point. When the GP is marginalized by taking the relevant submatrix block of the covariance matrix, this conditioning gives us our desired 1-dimensional prediction.

Also notice that the covariance does not depend on observations but scales quadratically to the norm of the testing inputs. This is a key feature of GPs. The variance is comprised of the difference between the prior covariance, $\mathbf{K}_{\mathbf{x}_* \mathbf{x}_*}$, and positive term $\mathbf{K}_{\mathbf{x}_* \mathbf{x}} [\mathbf{K}_{\mathbf{x} \mathbf{x}} + \sigma_n^2 \mathbb{1}_{n \times n}]^{-1} \mathbf{K}_{\mathbf{x}_* \mathbf{x}}^\top$ which represents knowledge given by the observations about the underlying function.

Algorithm 1 shows one possible implementation for computing the mean and covariance of a single

Algorithm 1: Unoptimized GPR

input : Observations \mathbf{X}, \mathbf{y} and a test input \mathbf{x}_* .

output: A prediction \bar{f}_* with its corresponding variance $\mathbb{V}[f_*]$.

- 1 $\mathbf{L} = \text{cholesky}(\mathbf{K}_{\mathbf{x} \mathbf{x}} + \sigma_n^2 \mathbb{1}_{n \times n})$
 - 2 $\boldsymbol{\alpha} = \text{lin-solve}(\mathbf{L}^\top, \text{lin-solve}(\mathbf{L}, \mathbf{y}))$
 - 3 $\bar{f}_* = \mathbf{K}_{\mathbf{x}_* \mathbf{x}} \boldsymbol{\alpha}$
 - 4 $\mathbf{v} = \text{lin-solve}(\mathbf{L}, \mathbf{K}_{\mathbf{x}_* \mathbf{x}})$
 - 5 $\mathbb{V}[f_*] = \mathbf{K}_{\mathbf{x}_* \mathbf{x}_*} - \mathbf{v}^\top \mathbf{v}$
 - 6 **return** $\bar{f}_*, \mathbb{V}[f_*]$
-

A Cholesky decomposition is typically used since \mathbf{L} can be used twice to assist in solving both the linear systems in the mean and covariance. Unfortunately, a Cholesky decomposition incurs a runtime of $\mathcal{O}(n^3)$ where n is the number of samples making it impractical for large data sets. In the later chapters we shall consider other methods for solving these linear systems.

0.6. Gaussian Processes for Classification. As with most classification models, the Gaussian processes classifier (GPC) seeks an estimate for the joint probability $p(y, \mathbf{x})$ where $\mathbf{x} \in \mathbb{R}^d$ is an input, as in the regression case, but y is now a class taking on a discrete and finite number of values $\{\mathcal{C}_i\}_{i=1}^C$. Using Baye's theorem the joint probability density can be decomposed into either $p(y)p(\mathbf{x} | y)$ or $p(\mathbf{x})p(y | \mathbf{x})$ giving rise to the *generative* and *discriminative* approaches respectively [Ras06, page 34]. The generative approach models the prior probabilities of each class, $p(\mathcal{C}_i)$, as well as the class conditional probabilities for each input $p(\mathbf{x} | \mathcal{C}_i)$ and computes the posterior as

$$p(y | \mathbf{x}) = \frac{p(y)p(\mathbf{x} | y)}{\sum_{i=1}^C p(\mathcal{C}_i)p(\mathbf{x} | \mathcal{C}_i)}.$$

On the other hand, the discriminative method focuses on modelling $p(y | \mathbf{x})$ directly. With both these paradigms at our disposal, which one would be preferred for our GPC? While there are strengths and weaknesses associated with both models, the discriminative approach is usually chosen as it has a rather attractive property of directly modeling what we require, that is $p(y | \mathbf{x})$. Additionally, the density estimation of $p(\mathbf{x} | \mathcal{C}_i)$ using in the generative model presents a number of difficulties, especially for larger values of d . If we are only focused on classifying inputs, the generative approach could mean trying to solve a harder problem than what is necessary. For this reason we focus on GPCs that adopt the discriminative approach.

0.6.1. *Linear Models for Classification.* We can start by reviewing linear models for the simplest form of classification, that is binary classification. Adopting the notation from SVM (see section 0.4.1) literature, the binary classification problem involves assigning an input \mathbf{x} to a class of either -1 or $+1$. For a linear model likelihood can be formulated as

$$(15) \quad p(y = +1 | \mathbf{x}, \mathbf{w}) = \sigma(\langle \mathbf{x}, \mathbf{w} \rangle)$$

given a weight vector \mathbf{w} and where $\sigma(z)$ is chosen to be any sigmoid function, see definition 14.

Definition 14 (Sigmoid Function). *A sigmoid function is a monotonically increasing function mapping from \mathbb{R} to $[0, 1]$ [Ras06, page 35].*

In this text, the commonly used logistic function

$$(16) \quad \sigma(z) = \frac{1}{1 + \exp(-z)}$$

will take the role of the sigmoid function in equation 15, graphed in Figure 5. This type of model is aptly named the logistic regression. Unlike GPR, the likelihood is no longer a Gaussian distribution. Instead

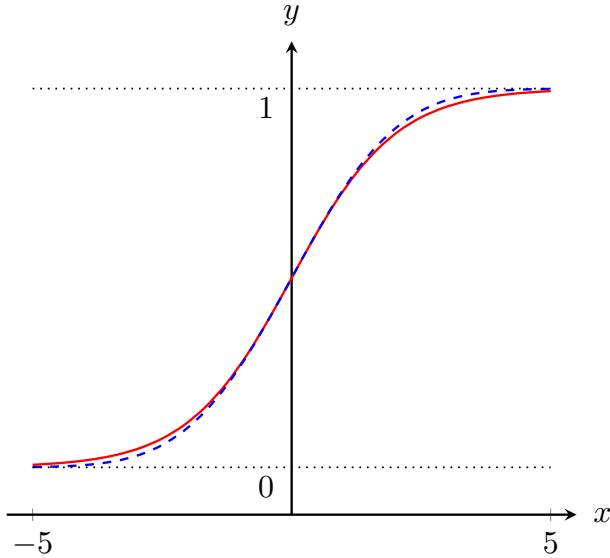


FIGURE 5. The logistic function from equation 16 (solid red) juxtaposed with a close approximation, the scaled probit function (dashed blue).

it follows the Bernoulli distribution

$$p(y | \mathbf{x}, \mathbf{w}) = \sigma(\langle \mathbf{x}, \mathbf{w} \rangle)^y (1 - \sigma(\langle \mathbf{x}, \mathbf{w} \rangle))^{1-y}$$

which for symmetric likelihood functions can be written more concisely as

$$p(y_i | \mathbf{x}_i, \mathbf{w}) = \sigma(y_i f_i)$$

where

$$(17) \quad f_i \triangleq f(\mathbf{x}_i) = \langle \mathbf{x}_i, \mathbf{w} \rangle.$$

Thus, the logistic regression model can be written as the log ratio of the likelihoods of the input belonging to either class, that is

$$\text{logit}(\mathbf{x}) \triangleq \langle \mathbf{x}, \mathbf{w} \rangle = \log \left(\frac{p(y=+1)}{p(y=-1)} \right)$$

where logit is commonly referred to as the logit transformation [Ras06, page 37]. For a given dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ we assume each observation is independently generated conditioned over $f(\mathbf{x})$. Similar to GPR, a Gaussian prior is used for the weights so that $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \sigma_p)$ giving an un-normalised log posterior of

$$\log p(\mathbf{w} | \mathbf{X}, \mathbf{y}) \propto -\frac{1}{2} \mathbf{w}^\top \Sigma_p^{-1} \mathbf{w} + \sum_{i=1}^n \log \sigma(y_i f_i).$$

However, unlike GPR an analytic form for the mean and variance for the posterior is not available due to the non-Gaussian nature of the likelihood, although, when using the logistic function it is easy enough to show that the log likelihood is concave as a function of \mathbf{w} for a fixed dataset. This means a number of numerical optimization techniques, such as Newton's method or the Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm [Fle00] can be used to solve these values.

The idea behind Gaussian process classification for binary classes is that a Gaussian process prior is placed over a latent function $f(\mathbf{x})$ where the output is then "squashed" through a sigmoid function to obtain a prior on

$$\pi(\mathbf{x}) \triangleq p(y=+1 | \mathbf{x}) = \sigma(f(\mathbf{x})).$$

This construction is illustrated in Figure 6 and provides a natural extension to the linear logistic regression model.

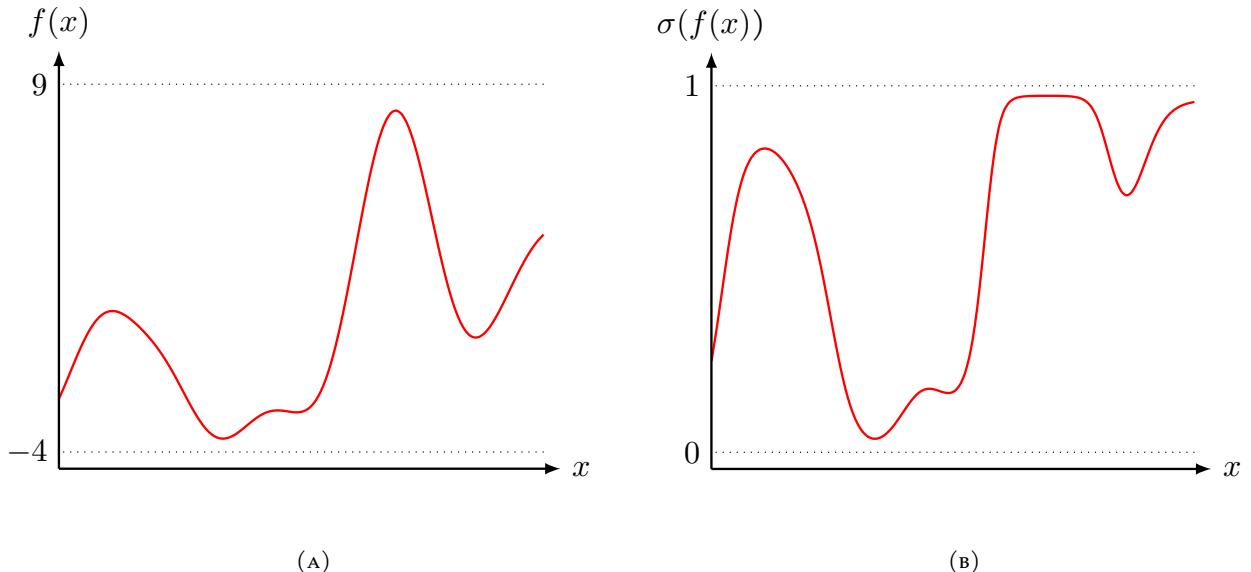


FIGURE 6. The latent function f , panel (A), is transformed using a sigmoid function, panel (B), to provide a probabilistic interpretation of x belonging to the class +1.

Specifically, the linear model from equation 17 is replaced with a GPR model and the Gaussian prior on the weights with a GPR weight prior with

$$p\left(\begin{bmatrix} \mathbf{f} \\ f_* \end{bmatrix}\right) = \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} \mathbf{K}_{XX} & \mathbf{K}_{x^*X}^\top \\ \mathbf{K}_{x^*X} & k(\mathbf{x}_*, \mathbf{x}_*) \end{bmatrix}\right)$$

where $f_* = f(\mathbf{x}_*)$ and $\mathbf{f} = f(\mathbf{X})$. For classification tasks, we assume that each observation has received the correct label which is why no noise is added to the covariance matrix.

Note that values of f are also never observed within the phenomena we are modelling, nor are we particularly interested in them. The function f serves the role of a *nuisance function* and acts solely as a convenience tool within our formulations. The ultimate goal is to make predictions for π , not f , and the goal of the coming sections will be to eventually integrate out f .

Subsequently, predictions for $\pi_* = \pi(\mathbf{x}_*)$ are made by average over all possible latent functions weighted by the posterior giving the prediction

$$(18) \quad \bar{\pi}_* \triangleq p(y_* = +1 | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) = \int \sigma(f_*) p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) df_*$$

While this is a sound model, computing predictions is not so straight forward since the integral in 18 is not analytically tractable for the same reason as the linear binary classifier. Later on we will see how we can make use of our numerical toolbox to derive a good approximation for $\bar{\pi}_*$.

0.6.2. Lapace Approximation for Posterior. We saw that the integral in 18 could not be used to make predictions for $\bar{\pi}_*$ analytically. In this section we shall address how the distribution for the latent process, $p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*)$, can be numerically approximated to provide a numerically tractable succedaneum. Using Baye's theorem

$$\begin{aligned} p(f_* | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) &= \int p(f_*, \mathbf{f} | \mathbf{X}, \mathbf{y}, \mathbf{x}_*) d\mathbf{f} \\ &= \frac{1}{p(\mathbf{y} | \mathbf{X}, \mathbf{x}_*)} \int p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) p(\mathbf{f} | \mathbf{X}) p(\mathbf{y} | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) d\mathbf{f} \\ &= \int p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) p(\mathbf{f} | \mathbf{X}, \mathbf{y}) d\mathbf{f} \end{aligned}$$

using the fact that $p(\mathbf{y} | \mathbf{X}, \mathbf{x}_*, \mathbf{f}, f_*) = p(\mathbf{y} | \mathbf{X}, \mathbf{x}_*, \mathbf{f})$ [Bis06, Ras06]. The conditional distribution $p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f})$ can be derived as

$$p(f_* | \mathbf{X}, \mathbf{x}_*, \mathbf{f}) = \mathcal{N}(f_* | \mathbf{K}_{x^*X} \mathbf{K}_{XX}^{-1} \mathbf{y}, k(\mathbf{x}_*, \mathbf{x}_*) - \mathbf{K}_{x^*X} \mathbf{K}_{XX}^{-1} \mathbf{K}_{x^*X}^\top)$$

through the use of equation 13 and 14. Unfortunately

$$p(\mathbf{f} | \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{f}) p(\mathbf{f} | \mathbf{X})}{p(\mathbf{y} | \mathbf{X})}$$

does not follow a Gaussian distribution. Instead we can use a Lapace approximation to estimate $p(\mathbf{f} | \mathbf{X}, \mathbf{y})$ as a Gaussian distribution. Briefly, the Lapace approximation works by assuming the distribution at hand, $p(\mathbf{z})$, can be modelled as

$$p(\mathbf{z}) = \frac{1}{c} q(\mathbf{z})$$

where $q(\mathbf{z})$ is multivariate Gaussian and c is some normalization constant [Bis06, page 214]. To do this, first the centre of $q(\mathbf{z})$ is placed at the mode of $p(\mathbf{z})$. The mode of $p(\mathbf{z})$ is

$$\mathbf{z}_0 = \arg \min_{\mathbf{z}} p(\mathbf{z})$$

which can be computed by solving

$$(19) \quad \nabla p(\mathbf{z}_0) = \mathbf{0}.$$

To ensure the covariance of the synthesized multivariate Gaussian behaves similar to the original distribution we can make use of an important property of the Gaussian distribution which is its logarithm being a quadratic function of its inputs. Taking the Taylor series expansion of $\ln q(\mathbf{z})$ centered at \mathbf{z}_0 yields

$$\ln q(\mathbf{z}) \simeq \ln q(\mathbf{z}_0) - \frac{1}{2} (\mathbf{z} - \mathbf{z}_0)^T \mathbf{A} (\mathbf{z} - \mathbf{z}_0)$$

where

$$\mathbf{A} = -\nabla \nabla \ln q(\mathbf{z})|_{\mathbf{z}=\mathbf{z}_0}.$$

Exponentiating both sides gives

$$(20) \quad \begin{aligned} q(\mathbf{z}) &\simeq q(\mathbf{z}_0) \exp \left(-\frac{1}{2} (\mathbf{z} - \mathbf{z}_0)^T \mathbf{A} (\mathbf{z} - \mathbf{z}_0) \right) \\ &\propto \mathcal{N}(\mathbf{z} | \mathbf{z}_0, \mathbf{A}^{-1}). \end{aligned}$$

Returning to our original problem of estimating $p(\mathbf{f} | \mathbf{X}, \mathbf{y}) \propto p(\mathbf{y} | \mathbf{f}) p(\mathbf{f} | \mathbf{X})$ as a Gaussian distribution, the prior $p(\mathbf{f} | \mathbf{X})$ follows a Gaussian distribution with zero mean and covariance $\mathbf{K}_{\mathbf{XX}}$ and the distribution of $p(\mathbf{y} | \mathbf{f})$ (assuming independence of samples) can be written as

$$p(\mathbf{y} | \mathbf{f}) = \prod_{i=1}^n \sigma(y_i f_i).$$

To find a Laplace approximation for $p(\mathbf{f} | \mathbf{X}, \mathbf{y})$ we only need to consider an unnormalized posterior when maximizing with respect to \mathbf{f} since $p(\mathbf{y} | \mathbf{f})$ does not depend on \mathbf{f} . Thus, the log of the unnormalized posterior is

$$\begin{aligned} \Psi(\mathbf{f}) &\triangleq \ln p(\mathbf{y} | \mathbf{f}) + \ln p(\mathbf{f} | \mathbf{X}) \\ &= - \sum_{i=1}^n \ln(1 + \exp(y_i f_i)) - \frac{1}{2} \mathbf{f}^T \mathbf{K}_{\mathbf{XX}}^{-1} \mathbf{f} - \frac{1}{2} \ln |\mathbf{K}_{\mathbf{XX}}| - \frac{n}{2} \ln 2\pi. \end{aligned}$$

The gradient and Hessian of the unnormalized posterior then becomes

$$\begin{aligned} \nabla \Psi(\mathbf{f}) &= \nabla \ln p(\mathbf{y} | \mathbf{f}) - \mathbf{K}_{\mathbf{XX}}^{-1} \mathbf{f} = (\mathbf{t} - \boldsymbol{\pi}) - \mathbf{K}_{\mathbf{XX}}^{-1} \mathbf{f} \\ \nabla \nabla \Psi(\mathbf{f}) &= \nabla \nabla \ln p(\mathbf{y} | \mathbf{f}) - \mathbf{K}_{\mathbf{XX}}^{-1} = -\mathbf{W} - \mathbf{K}_{\mathbf{XX}}^{-1} \end{aligned}$$

where $\pi_i = p(y_i = +1 | f_i) = \sigma(f_i)$, $\mathbf{t} = (\mathbf{y} + \mathbf{1}) / 2 \in \mathbb{R}^n$ and $\mathbf{W} \triangleq -\nabla \nabla \ln p(\mathbf{y} | \mathbf{f})$ is a diagonal matrix (since the distribution of y_i only depends on f_i and not $f_{j \neq i}$) with entries $\mathbf{W}_{ii} = \sigma(y_i f_i)$ [Bis06, Ras06]. From equation 19, the mode of $\hat{\mathbf{f}}$ of Ψ can be computed as

$$(21) \quad \begin{aligned} \nabla \Psi(\hat{\mathbf{f}}) &= \mathbf{0} = (\mathbf{t} - \boldsymbol{\pi}) - \mathbf{K}_{\mathbf{XX}}^{-1} \hat{\mathbf{f}} \\ \iff \hat{\mathbf{f}} &= \mathbf{K}_{\mathbf{XX}} (\mathbf{t} - \boldsymbol{\pi}). \end{aligned}$$

Since $\mathbf{t} - \boldsymbol{\pi}$ is a non-linear function, a non-linear optimization technique method is required to solve $\hat{\mathbf{f}}$ in 21. Since the Hessian of $\Psi(\mathbf{f})$ is available, Newton's method is typically employed as fast iterative method to approximate $\hat{\mathbf{f}}$ where $\hat{\mathbf{f}}$ is updated as

$$\hat{\mathbf{f}}^{\text{new}} = \mathbf{K}_{\mathbf{XX}} (\mathbb{1}_{n \times n} + \mathbf{W} \mathbf{K}_{\mathbf{XX}})^{-1} \left(\mathbf{W} \hat{\mathbf{f}}^{\text{old}} + \nabla \ln (\mathbf{y} | \hat{\mathbf{f}}^{\text{old}}) \right).$$

Once a suitable mode is found, using equation 20, the Lapacian approximation for $p(\mathbf{f} | \mathbf{X}, \mathbf{y})$ becomes

$$(22) \quad p(\mathbf{f} | \mathbf{X}, \mathbf{y}) \simeq q(\mathbf{f} | \mathbf{X}, \mathbf{y}) = \mathcal{N} \left(\hat{\mathbf{f}}, (\mathbf{K}_{\mathbf{XX}}^{-1} + \mathbf{W})^{-1} \right).$$

0.6.3. *Predictions.* With the Lapace approximation for $p(\mathbf{f} | \mathbf{X}, \mathbf{y})$ (equation 22) and an exact probability distribution for $p(f_{\star} | \mathbf{X}, \mathbf{x}_{\star}, \mathbf{f})$, a mean for the latent process, $p(f_{\star} | \mathbf{X}, \mathbf{y}, \mathbf{x}_{\star})$, can now be computed by invoking 13 to give

$$(23) \quad \begin{aligned} \mu_{f_{\star}} &= \mathbb{E}[f_{\star} | \mathbf{X}, \mathbf{y}, \mathbf{x}_{\star}] = \mathbf{K}_{\mathbf{x}_{\star} \mathbf{X}} \mathbf{K}_{\mathbf{XX}}^{-1} \hat{\mathbf{f}} \\ &= \mathbf{K}_{\mathbf{x}_{\star} \mathbf{X}} \nabla \ln (\mathbf{y} | \hat{\mathbf{f}}) \\ &= \mathbf{K}_{\mathbf{x}_{\star} \mathbf{X}} (\mathbf{t} - \boldsymbol{\pi}). \end{aligned}$$

Similarly, the variance can be computed using equation 14 to give

$$(24) \quad \sigma_{f_{\star}}^2 = \mathbb{V}[f_{\star} | \mathbf{X}, \mathbf{y}, \mathbf{x}_{\star}] = k(\mathbf{x}_{\star}, \mathbf{x}_{\star}) - \mathbf{K}_{\mathbf{x}_{\star} \mathbf{X}} (\mathbf{K}_{\mathbf{XX}} + \mathbf{W}^{-1})^{-1} \mathbf{K}_{\mathbf{x}_{\star} \mathbf{X}}^{\top}.$$

Using equation 18, predictions can now be made as

$$(25) \quad \bar{\pi}_{\star} \simeq \int \sigma(f_{\star}) q(f_{\star} | \mathbf{X}, \mathbf{y}, \mathbf{x}_{\star}) df_{\star}$$

where $q(f_{\star} | \mathbf{X}, \mathbf{y}, \mathbf{x}_{\star})$ is a multivariate Gaussian distribution with mean and variance given by equations 23 and 24 respectively. Notice that the prediction given in 25 is a convolution of a Gaussian and logistic function which unfortunately cannot be evaluated analytically. However, Spiegelhalter and Lauritzen [Spi90] show that a good approximation can be found by replacing the sigmoid function with the probit function $\Phi(\lambda a)$ which is simply the cumulative distribution function (CDF) of the standard Gaussian distribution. To get the best approximation using the probit function, the constant factor λ is adjusted to equate their slopes at the origin. The value of λ that gives this equality is $\lambda = \sqrt{\pi/8}$. The similarity between the sigmoid function and probit function rescaled by a factor of $\sqrt{\pi/8}$ is illustrated in Figure 5. The reason for replacing the sigmoid function with a probit function is that the convolution of a Gaussian distribution and probit function can be analytically evaluated as

$$(26) \quad \int \Phi(\lambda a) \mathcal{N}(a | \mu, \sigma^2) da = \Phi \left(\frac{\mu}{(\lambda^{-2} + \sigma^2)^{\frac{1}{2}}} \right).$$

Again apply the approximation $\sigma(a) \simeq \Phi(\lambda a)$ to left hand side of 26 gives the following estimate for the convolution of a Gaussian and sigmoid function

$$(27) \quad \int \sigma(a) \mathcal{N}(a | \mu, \sigma^2) da \simeq \sigma \left(\frac{\mu}{(1 + \pi\sigma^2/8)^{\frac{1}{2}}} \right)$$

[Bis06, page 219]. The integral used to approximate $\bar{\pi}_*$ in 25 can now be estimated using 27 to give

$$\bar{\pi}_* = \sigma \left(\frac{\mu_{f_*}}{\left(1 + \pi\sigma_{f_*}^2/8\right)^{1/2}} \right).$$

This theory justifies Algorithm 2 which creates predictions based on the GPC method.

Algorithm 2: Unoptimized GPC

input : Observations \mathbf{X}, \mathbf{y} and a test input \mathbf{x}^* .
output: A prediction \bar{f}_* with its corresponding variance $\mathbb{V}[f_*]$.

```

1  $t = (\mathbf{y} + \mathbf{1}) / 2$ 
2  $f = \mathbf{0}$ 
3 repeat
4    $W = \text{diag}(\sigma(\mathbf{y}.^*\mathbf{f}))$ 
5    $\alpha = \text{lin-solve}(\mathbf{1}_{n \times n} + W\mathbf{K}_{XX}, \mathbf{K}_{XX})$ 
6    $f = \alpha(t - \sigma(f) + Wf)$ 
7 until convergence
8  $\mu_{f_*} = \mathbf{K}_{x_*\mathbf{X}}(t - \sigma(f))$ 
9  $\sigma_{f_*}^2 = k(x_*, x_*) - \mathbf{K}_{x_*\mathbf{X}}(\mathbf{K}_{XX} + W^{-1})^{-1}\mathbf{K}_{x_*\mathbf{X}}^\top$ 
10  $\bar{\pi}_* = \sigma \left( \mu_{f_*} / \left(1 + \pi\sigma_{f_*}^2/8\right)^{1/2} \right)$ 
11 return  $\bar{\pi}_*, \mu_{f_*}, \sigma_{f_*}^2$ 
```

1. APPLICATIONS AND RESULTS

So far we have discussed a number of different approximation techniques to use to speed up some of the bottle necks within our Gaussian Process regression and classification algorithms. Here, we shall see how they can be used within the naive implementations from ?? to provide implementations that produce much faster predictions, at the expense of slightly less accurate results. We shall also go through an extensive treatment of how the various methods reviewed in previous chapters were implemented and the setup of the experiments used to make comparisons to determine the most successful method.

1.1. Gaussian Processes Prediction Reviewed. In Chapter , a naive implementation for Gaussian Process prediction was presented in Algorithm 1. While this does provide a simple and convenient way to produce predictions for a regression task, it is not quite all smooth sailing from here. Unfortunately, there are a number of problems with this algorithm in terms of scalability. For convenience, this algorithm has been restated below but this time with its various bottleneck highlighted in red.

Algorithm 1: Unoptimized GPR

input : Observations \mathbf{X}, \mathbf{y} and a test input \mathbf{x}_* .
output: A prediction \bar{f}_* with its corresponding variance $\mathbb{V}[f_*]$.

```

1 L = cholesky ( $\mathbf{K}_{\mathbf{XX}} + \sigma_n^2 \mathbf{1}_{n \times n}$ )
2 α = lin-solve ( $\mathbf{L}^\top$ , lin-solve ( $\mathbf{L}, \mathbf{y}$ ))
3  $\bar{f}_* = \mathbf{K}_{\mathbf{x}_* \mathbf{X}} \boldsymbol{\alpha}$ 
4 v = lin-solve ( $\mathbf{L}, \mathbf{K}_{\mathbf{x}_* \mathbf{X}}$ )
5  $\mathbb{V}[f_*] = \mathbf{K}_{\mathbf{x}_* \mathbf{x}_*} - \mathbf{v}^\top \mathbf{v}$ 
6 return  $\bar{f}_*, \mathbb{V}[f_*]$ 
```

To start computing the kernel matrix, $\mathbf{K}_{\mathbf{XX}}$, on line 1 carries an $\mathcal{O}(n^2)$ runtime since $\mathcal{O}(n^2)$ pairwise kernel evaluations must be made. Moreover, solving linear systems using a cholesky decomposition (seen on lines 1, 2 and 4) will incur a runtime of $\mathcal{O}(n^3)$. While this quadratic scaling may not be much of an issue for smaller datasets, however this will scale very poorly for larger datasets making and training sets larger than 10^5 infeasible to use on most modern work stations [WJMaSKaAGW21, page 2]. The Gaussian Process classifier from Algorithm 2 suffers for the exact same reasons as GPR. To mitigate this computation burden of these bottle necks, we can replace these procedures with their approximative counterparts discussed in Chapters ??,?? and ??.

To start, computing the kernel matrix can be done using either the Nystrom method or the RFF technique, both of which provide better asymptotic runtimes. Similarly, the Cholesky decomposition and linear solvers can be replaced with either CG or MINRES to, again, improve runtime performance. This especially makes sense approximating the kernel matrix as the Cholesky decomposition will provide an almost-exact solution to solving these linear system, although this is rather counterproductive if approximations are used earlier on within the algorithm.

1.2. Experimental Setup. Each method was implemented in python3 and run using the python3.8.5 interrupter distributed by the official python website.

REFERENCES

- [Ras06] Carl Edward and Williams Rasmussen Christopher K. I, *Gaussian processes for machine learning / Carl Edward Rasmussen, Christopher K.I. Williams.*, Adaptive computation and machine learning, MIT Press, Cambridge, Mass., 2006 (eng).
- [HHF73] H. Howard Frisinger, *Aristotle's legacy in meteorology*, Bulletin of the American Meteorological Society **54** (1973), no. 3, 198–204.
- [Yul27] G. Udny Yule, *On a Method of Investigating Periodicities in Disturbed Series, with Special Reference to Wolfer's Sunspot Numbers*, Philosophical transactions of the Royal Society of London. Series A, Containing papers of a mathematical or physical character **226** (1927), no. 636-646, 267–298 (eng).
- [Box08] George E. P. and Jenkins Box Gwilym M and Reinsel, *Time series analysis : forecasting and control / George E.P. Box, Gwilym M. Jenkins, Gregory C. Reinsel.*, 4th ed., Wiley series in probability and statistics, John Wiley, Hoboken, N.J., 2008 (eng).
- [VdW19] Mark Van der Wilk, *Sparse Gaussian process approximations and applications*, University of Cambridge, 2019.
- [Cao18] Yanshuai Cao, *Scaling Gaussian Processes*, University of Toronto (Canada), 2018.
- [SD22] Matías and Estévez Salinero-Delgado José and Pipia, *Monitoring Cropland Phenology on Google Earth Engine Using Gaussian Process Regression*, Remote Sensing **14** (2022), no. 1, DOI 10.3390/rs14010146.
- [Pot13] Andries and Lawson Potgieter Kenton and Huete, *Determining crop acreage estimates for specific winter crops using shape attributes from sequential MODIS imagery*, International Journal of Applied Earth Observation and Geoinformation **23** (2013), DOI 10.1016/j.jag.2012.09.009.
- [NdF13] Nando de Freitas, *University of British Columbia CPSC 540, Lecture notes in Machine Learning*, University of British Columbia, 2013.
- [Mur12] Kevin P. Murphy, *Machine learning : a probabilistic perspective / Kevin P. Murphy.*, Adaptive computation and machine learning, MIT Press, Cambridge, MA, 2012 (eng).
- [Ber96] Z.G. Sheftel Berezansky G.F, *Functional analysis. Volume 1 / Y.M. Berezansky, Z.G. Sheftel, G.F. Us ; translated from the Russian by Peter V. Malyshov.*, 1st ed. 1996., Operator

- Theory: Advances and Applications, 85, Basel ; Boston ; Berlin : BirkhaIuser Verlag, Basel ; Boston ; Berlin, 1996 (eng).
- [Tre97] Lloyd N. (Lloyd Nicholas) and Bau Trefethen David, *Numerical linear algebra / Lloyd N. Trefethen, David Bau.*, SIAM Society for Industrial and Applied Mathematics, Philadelphia, 1997 (eng).
- [Dem97] James W Demmel, *Applied numerical linear algebra / James W. Demmel.*, Society for Industrial and Applied Mathematics, Philadelphia, Pa., 1997 (eng).
- [Ste08] Ingo and Christmann Steinwart Andreas, *Support Vector Machines*, 1st ed. 2008., Information Science and Statistics, Springer New York, New York, NY, 2008 (eng).
- [Ber03] Alain and Thomas-Agnan Berlinet Christine, *Reproducing Kernel Hilbert Spaces in Probability and Statistics*, Springer, SpringerLink (Online service), Boston, MA, 2003 (eng).
- [Ste99] Michael L Stein, *Interpolation of Spatial Data Some Theory for Kriging / by Michael L. Stein.*, 1st ed. 1999., Springer Series in Statistics, Springer New York : Imprint: Springer, New York, NY, 1999 (eng).
- [Bos92] Bernhard and Guyon Boser Isabelle and Vapnik, *A training algorithm for optimal margin classifiers*, Proceedings of the fifth annual workshop on computational learning theory, 1992, pp. 144–152 (eng).
- [Cor95] Corinna Cortes, *Support-Vector Networks*, Machine learning **20** (1995), no. 3, 273 (eng).
- [Kro14] Dirk P and C.C. Chan Kroese Joshua, *Statistical Modeling and Computation by Dirk P. Kroese, Joshua C.C. Chan.*, 1st ed. 2014., Springer New York : Imprint: Springer, New York, NY, 2014 (eng).
- [Fle00] R Fletcher, *Practical Methods of Optimization*, John Wiley and Sons, Incorporated, New York, 2000 (eng).
- [Bis06] Christopher M Bishop, *Pattern recognition and machine learning / Christopher M. Bishop.*, Information science and statistics, Springer, New York, 2006 (eng).
- [Spi90] David J and Lauritzen Spiegelhalter Steffen L, *Sequential updating of conditional probabilities on directed graphical structures*, Networks **20** (1990), no. 5, 579–605.
- [WJMaSKaAGW21] Wesley J. Maddox and Sanyam Kapoor and Andrew Gordon Wilson, *When are Iterative Gaussian Processes Reliably Accurate?*, CoRR **abs/2112.15246** (2021), available at [2112.15246](https://arxiv.org/abs/2112.15246).

- [MWM11] Michael W. Mahoney, *Randomized algorithms for matrices and data*, CoRR **abs/1104.5557** (2011).
- [Hal11] Nathan and Martinsson Halko Per-Gunnar and Tropp, *Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions*, SIAM review **53** (2011), no. 2, 217–288.
- [PGMaJT21] Per-Gunnar Martinsson and Joel Tropp, *Randomized Numerical Linear Algebra: Foundations and Algorithms*, arXiv, 2021.
- [FR20] Fred Roosta, *University of Queensland MATH3204, Lecture notes in Numerical Linear Algebra and Optimisation*, University of Queensland, 2020.
- [Dri06] Petros and Kannan Drineas Ravi and Mahoney, *Fast Monte Carlo Algorithms for Matrices I: Approximating Matrix Multiplication*, SIAM Journal on Computing **36** (2006), no. 1, 132-157, DOI 10.1137/S0097539704442684, available at <https://doi.org/10.1137/S0097539704442684>.
- [PDaMWM17] Petros Drineas and Michael W. Mahoney, *Lectures on Randomized Numerical Linear Algebra*, arXiv, 2017.
- [PDaMWM05] Petros Drineas and Michael W. Mahoney, *On the Nystrom Method for Approximating a Gram Matrix for Improved Kernel-Based Learning*, Journal of Machine Learning Research **6** (2005), no. 72, 2153-2175.
- [AGaMWM13] Alex Gittens and Michael W. Mahoney, *Revisiting the Nystrom Method for Improved Large-Scale Machine Learning*, CoRR **abs/1303.1849** (2013), available at [1303.1849](https://doi.org/10.4236/jmlr.v14i1.3430).
- [CMaCM17] Cameron Musco and Christopher Musco, *Recursive Sampling for the Nystrom Method*, arXiv, 2017.
- [PDe11] Petros Drineas etal., *Fast approximation of matrix coherence and statistical leverage*, CoRR **abs/1109.3843** (2011).
- [MBCaCMaCM15] Michael B. Cohen and Cameron Musco and Christopher Musco, *Ridge Leverage Scores for Low-Rank Approximation*, CoRR **abs/1511.07263** (2015).
- [Kum09] Sanjiv and Mohri Kumar Mehryar and Talwalkar, *Sampling techniques for the nystrom method*, Artificial intelligence and statistics, 2009, pp. 304–311.
- [Hoa78] David C and Welsch Hoaglin Roy E, *The Hat Matrix in Regression and ANOVA*, The American statistician **32** (1978), no. 1, 17–22 (eng).

- [Ala15] Ahmed and Mahoney Alaoui Michael W, *Fast Randomized Kernel Ridge Regression with Statistical Guarantees*, Advances in Neural Information Processing Systems, 2015.
- [Pre92] William H. (William Henry) Press, *Numerical recipes in C : the art of scientific computing / William H. Press ... [et al.]*, 2nd ed., Cambridge University Press, Cambridge, 1992 (eng).
- [Wan] Guorong and Wei Wang Yimin and Qiao, *Generalized Inverses: Theory and Computations*, Developments in Mathematics, vol. 53, Springer Singapore, Singapore (eng).
- [Gre97] Anne Greenbaum, *Iterative methods for solving linear systems* Anne Greenbaum., Frontiers in applied mathematics ; 17, Society for Industrial and Applied Mathematics SIAM, 3600 Market Street, Floor 6, Philadelphia, PA 19104, Philadelphia, Pa., 1997 (eng).
- [Cho07] Sou-Cheng (Terrya) Choi, *Iterative methods for singular linear equations and least - squares problems*, ProQuest Dissertations Publishing, 2007 (eng).
- [CHO11] Sou-Cheng T and PAIGE CHOI Christopher C and SAUNDERS, *MINRES-QLP: A KRYLOV SUBSPACE METHOD FOR INDEFINITE OR SINGULAR SYMMETRIC SYSTEMS*, SIAM journal on scientific computing **33** (2011), no. 3-4, 1810–1836 (eng).
- [Rah08] Ali and Recht Rahimi Benjamin, *Random Features for Large-Scale Kernel Machines*, Advances in Neural Information Processing Systems, 2008.
- [Pot21] Andres and Wu Potapczynski Luhuan and Biderman, *Bias-Free Scalable Gaussian Processes via Randomized Truncations* (2021) (eng).
- [Hah33] Hans Hahn, *S. Bochner, Vorlesungen über Fouriersche Integrale: Mathematik und ihre Anwendungen, Bd. 12.) Akad. Verlagsges., Leipzig 1932, VIII. u. 229S. Preis brosch. RM 14,40, geb. RM16*, Monatshefte für Mathematik **40** (1933), no. 1, A27–A27 (ger).
- [Liu21] Fanghui and Huang Liu Xiaolin and Chen, *Random Features for Kernel Approximation: A Survey on Algorithms, Theory, and Beyond*, IEEE transactions on pattern analysis and machine intelligence **PP** (2021) (eng).
- [HAe16] Haim Avron etal, *Quasi-Monte Carlo Feature Maps for Shift-Invariant Kernels*, Journal of Machine Learning Research **17** (2016), no. 120, 1-38.
- [DJSaJS15] Danica J. Sutherland and Jeff Schneider, *On the Error of Random Fourier Features*, 2015.

- [Yu16] Felix X and Suresh Yu Ananda Theertha and Choromanski, *Orthogonal Random Features* (2016) (eng).
- [Bro91] Peter J and Davis Brockwell Richard A, *Time Series: Theory and Methods*, Second Edition., Springer Series in Statistics, Springer New York, SpringerLink (Online service), New York, NY, 1991 (eng).
- [Cho17] Krzysztof and Rowland Choromanski Mark and Weller, *The Unreasonable Effectiveness of Structured Random Orthogonal Embeddings* (2017) (eng).
- [FaA76] Fino and Algazi, *Unified Matrix Treatment of the Fast Walsh-Hadamard Transform*, IEEE transactions on computers **C-25** (1976), no. 11, 1142–1146 (eng).
- [And15] Alexandr and Indyk Andoni Piotr and Laarhoven, *Practical and Optimal LSH for Angular Distance* (2015) (eng).
- [Cho20] Krzysztof and Likhoshevstov Choromanski Valerii and Dohan, *Rethinking Attention with Performers* (2020) (eng).
- [Boj16] Mariusz and Choromanska Bojarski Anna and Choromanski, *Structured adaptive and random spinners for fast machine learning computations* (2016) (eng).

APPENDIX A. SUPPLEMENTARY RESULTS

Additional results that may have not been included in the main text for conciseness.

A.1. Gram Matrix Spectral Values.

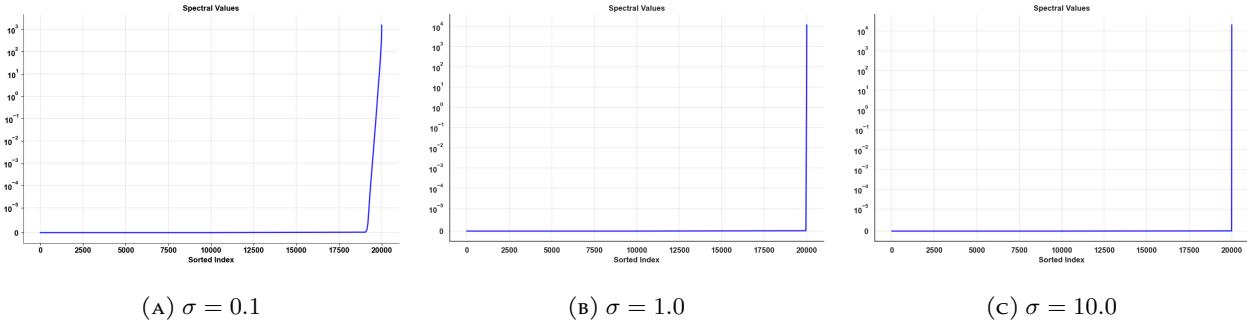


FIGURE 7. Spectral values for 3D-spatial network data.

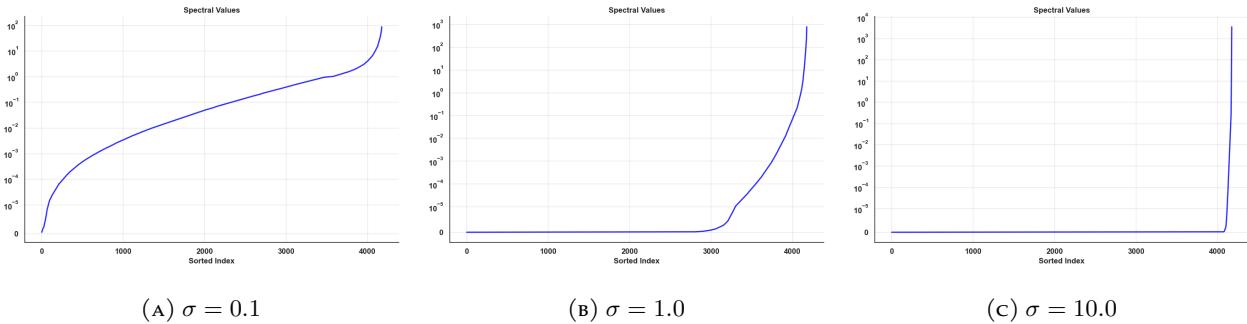


FIGURE 8. Spectral values for abalone data.

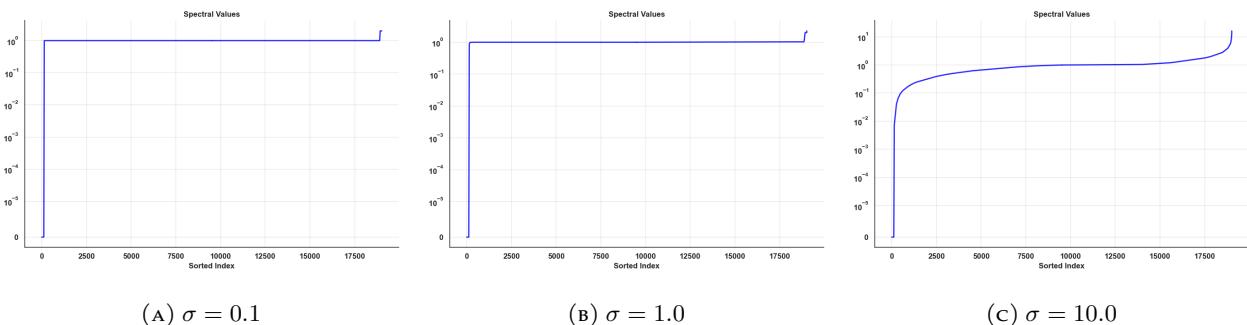


FIGURE 9. Spectral values for magic04 data.

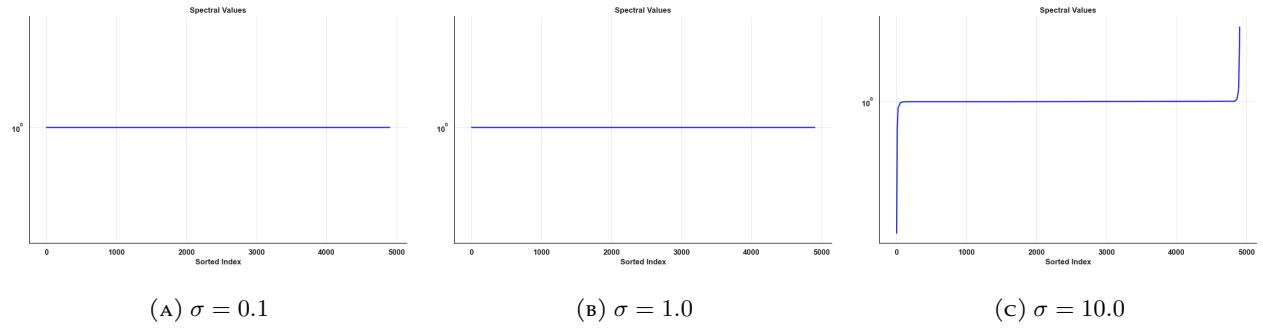


FIGURE 10. Spectral values for stock market data.

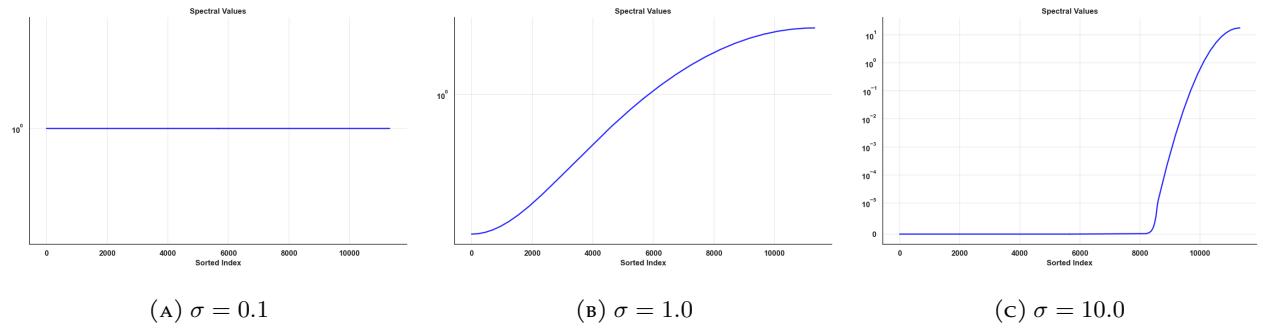


FIGURE 11. Spectral values for temperature data.

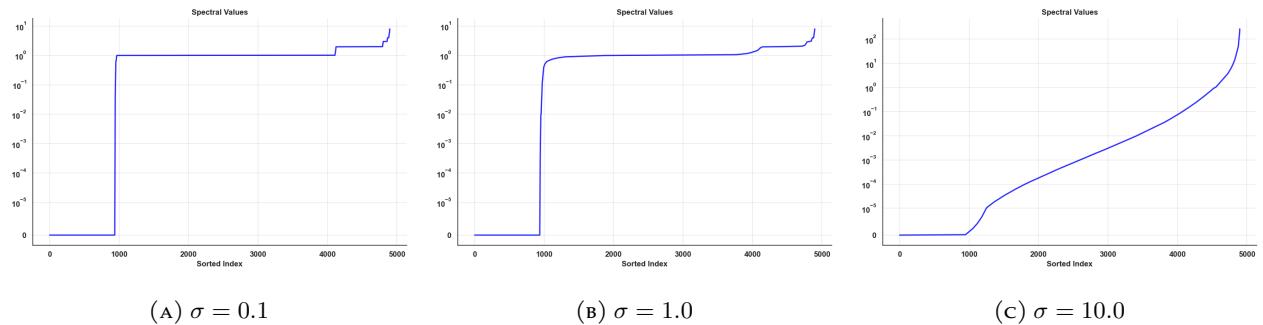


FIGURE 12. Spectral values for wine data.

A.2. Nystrom Scores.

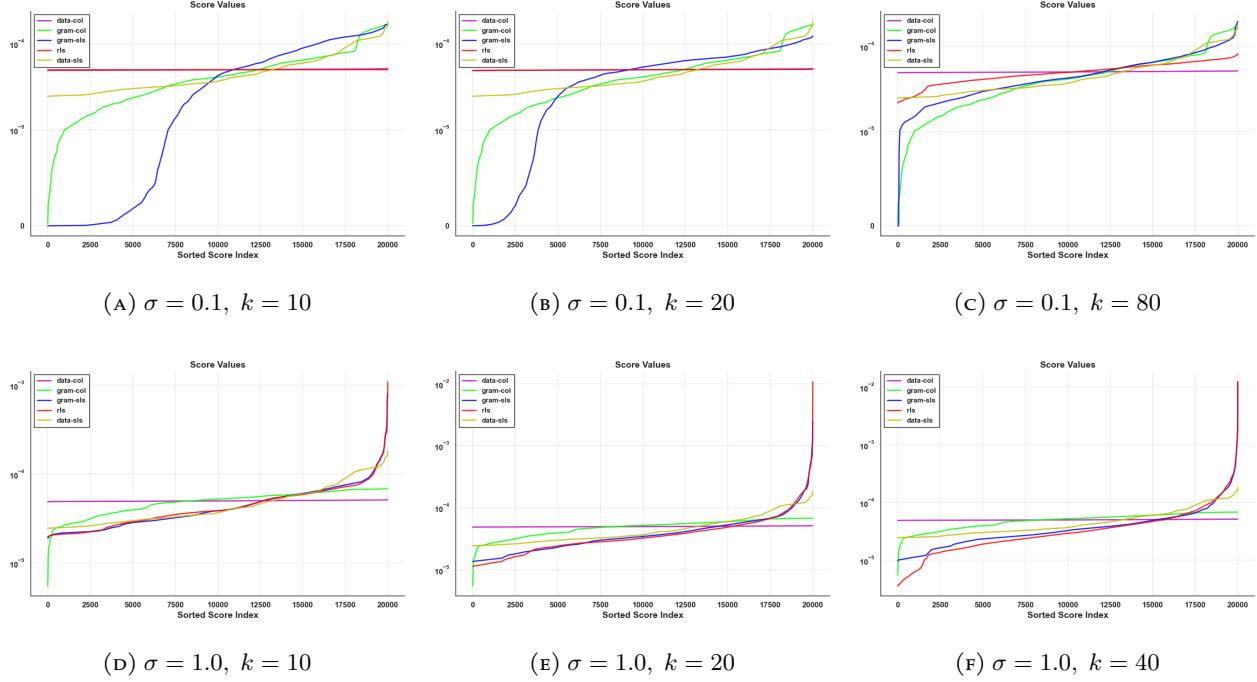


FIGURE 13. Nystrom scores for the 3D-spatial network data set.

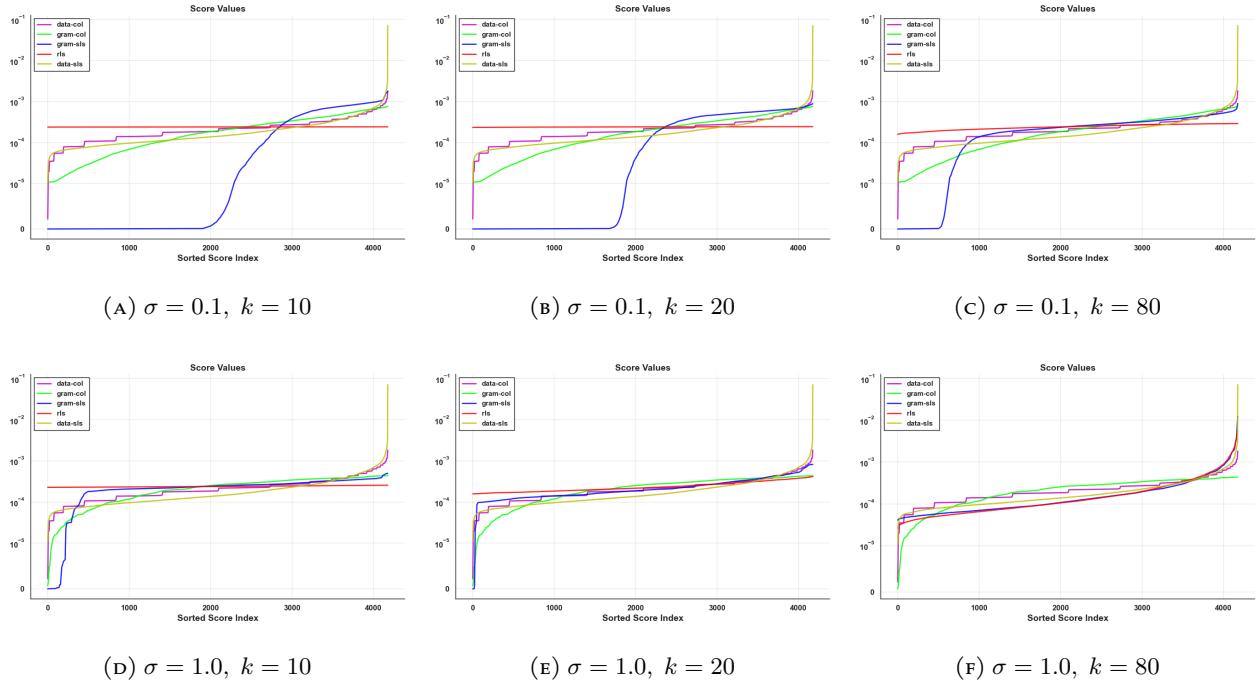


FIGURE 14. Nystrom scores for the Abalone data set.

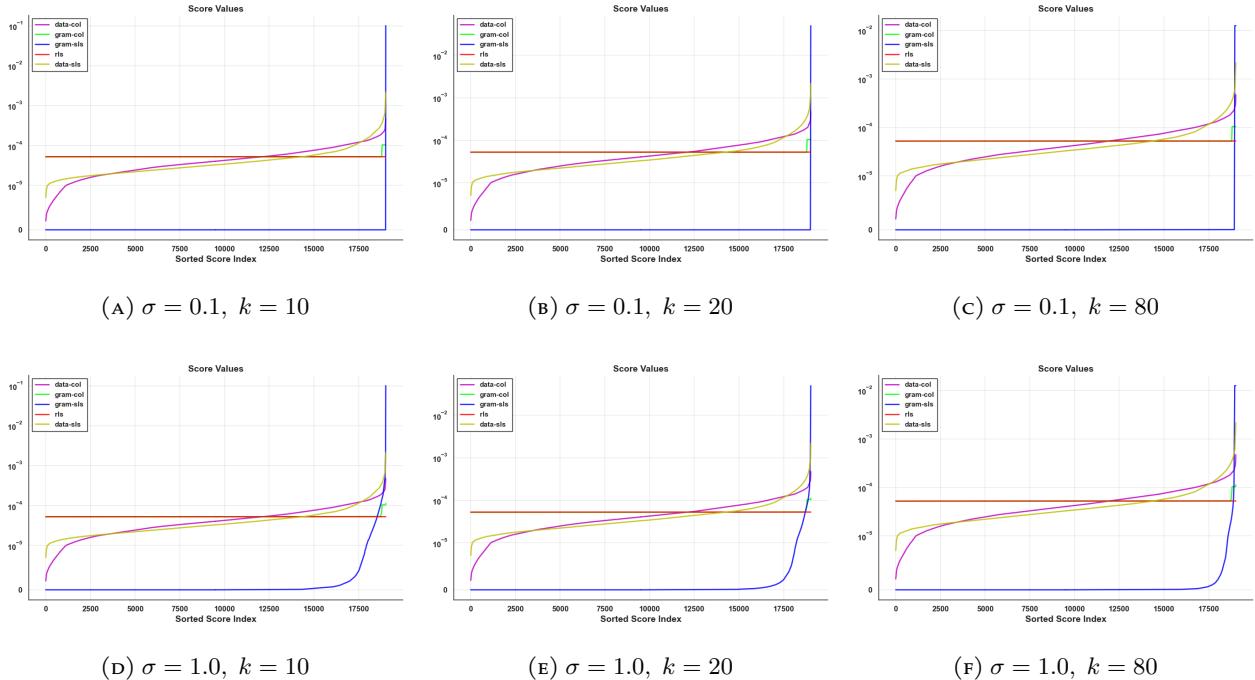


FIGURE 15. Nystrom scores for the Magic data set.

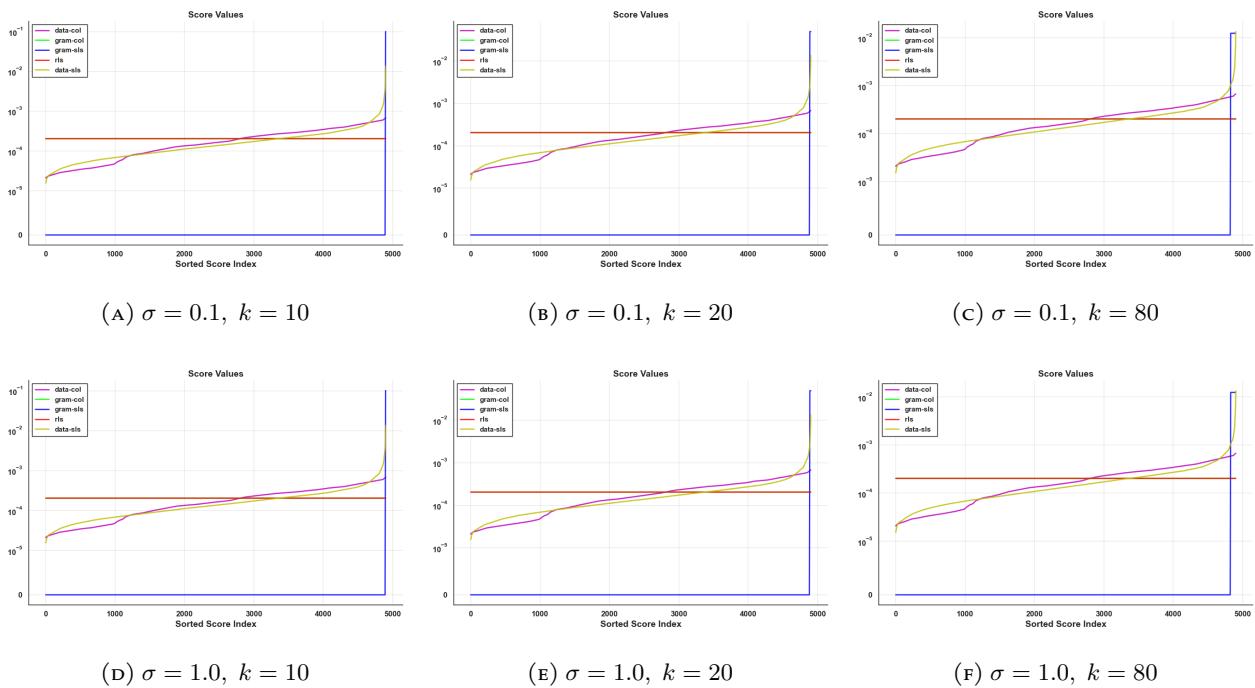


FIGURE 16. Nystrom scores for the stock market data set.

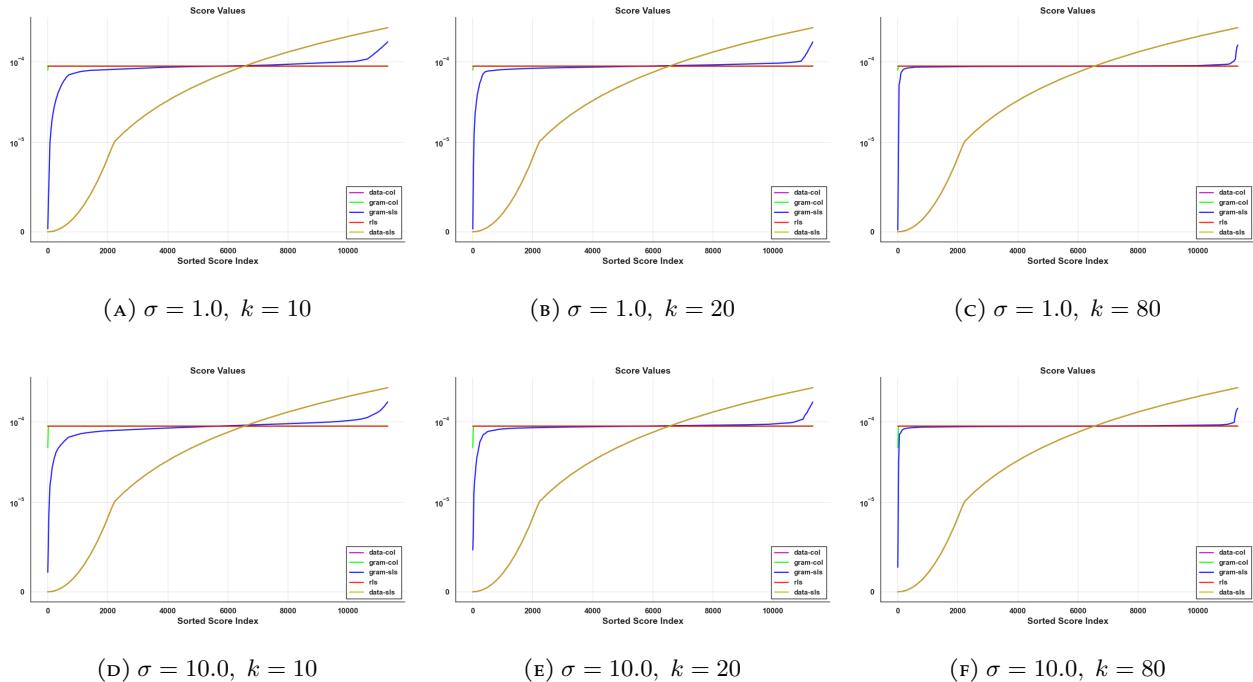


FIGURE 17. Nystrom scores for the temperature data set.

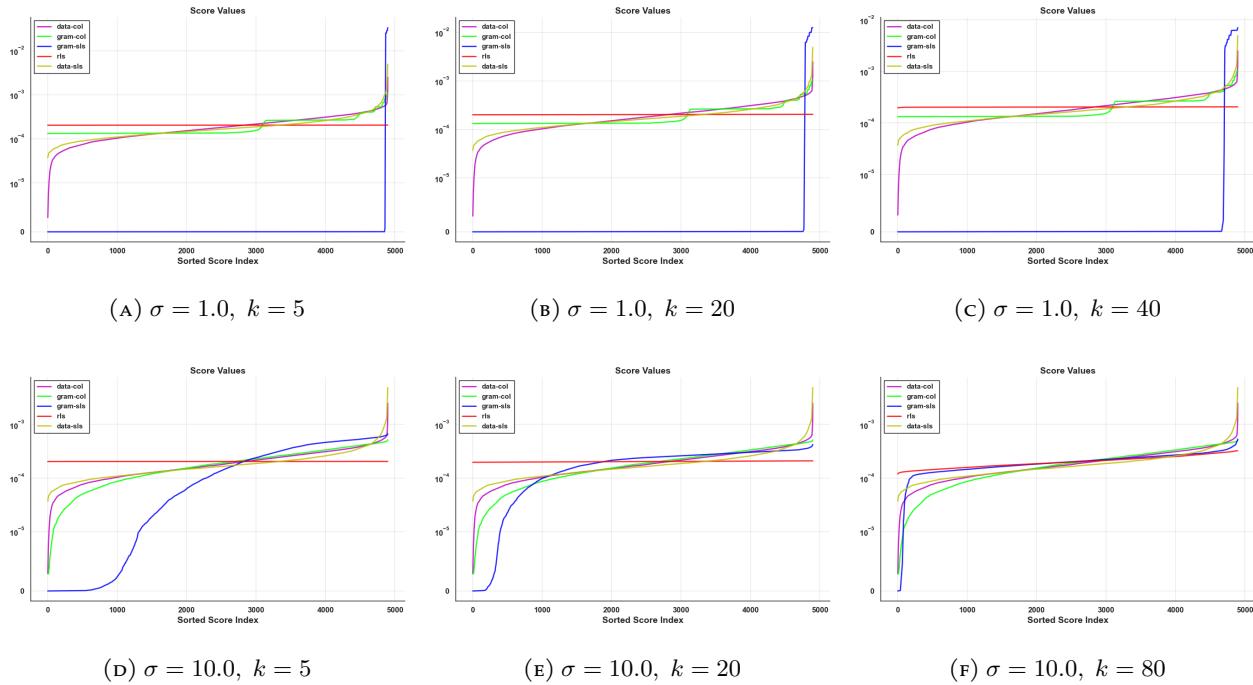


FIGURE 18. Nystrom scores for the wine data set.

A.3. Ridge Leverage Scores.

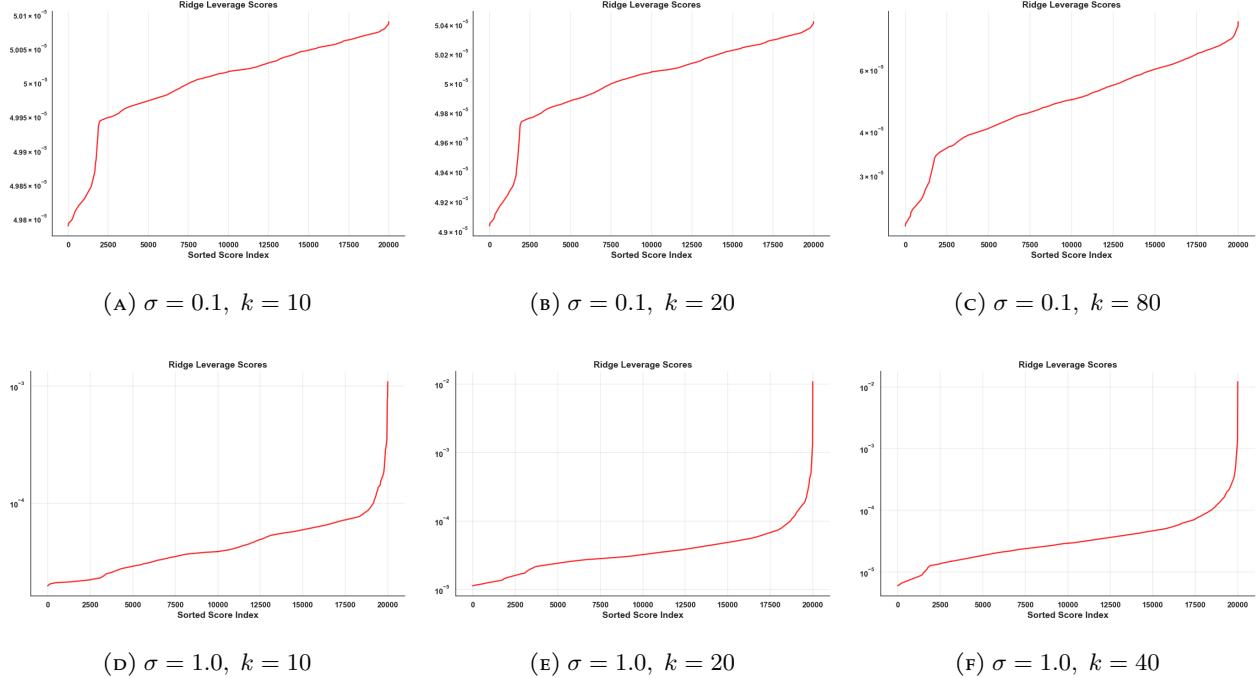


FIGURE 19. Ridge Leverage scores for the 3D-spatial network data set.

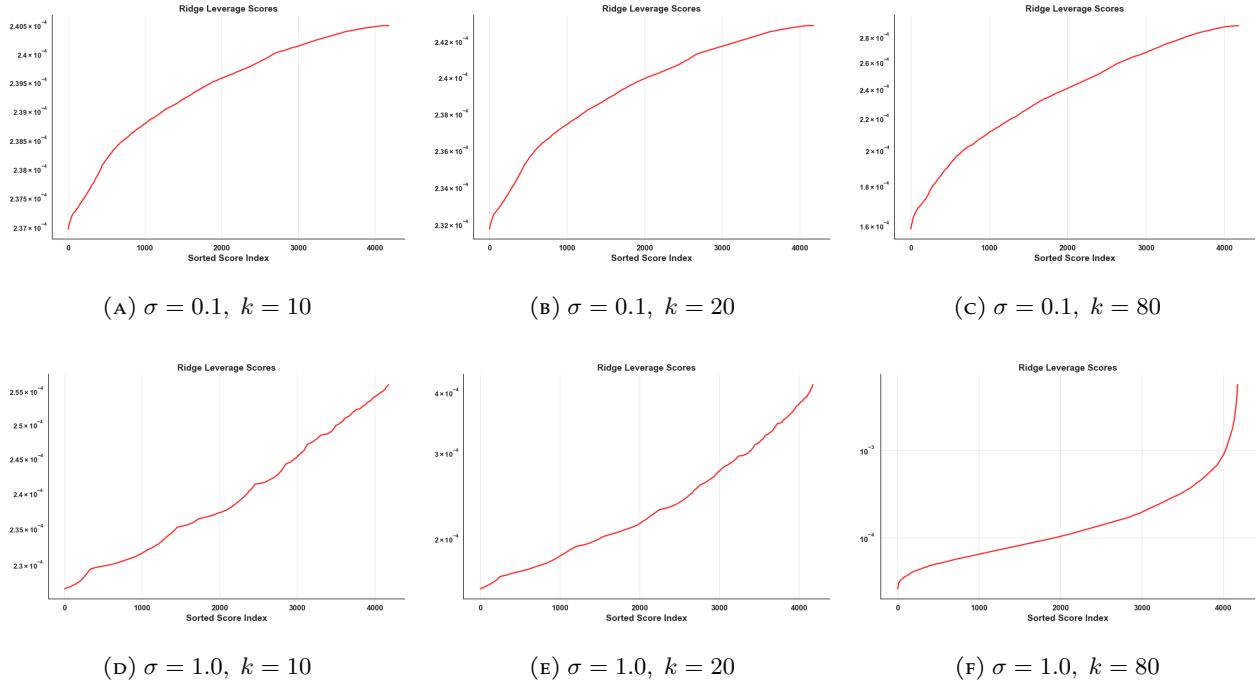


FIGURE 20. Ridge Leverage scores for the Abalone data set.

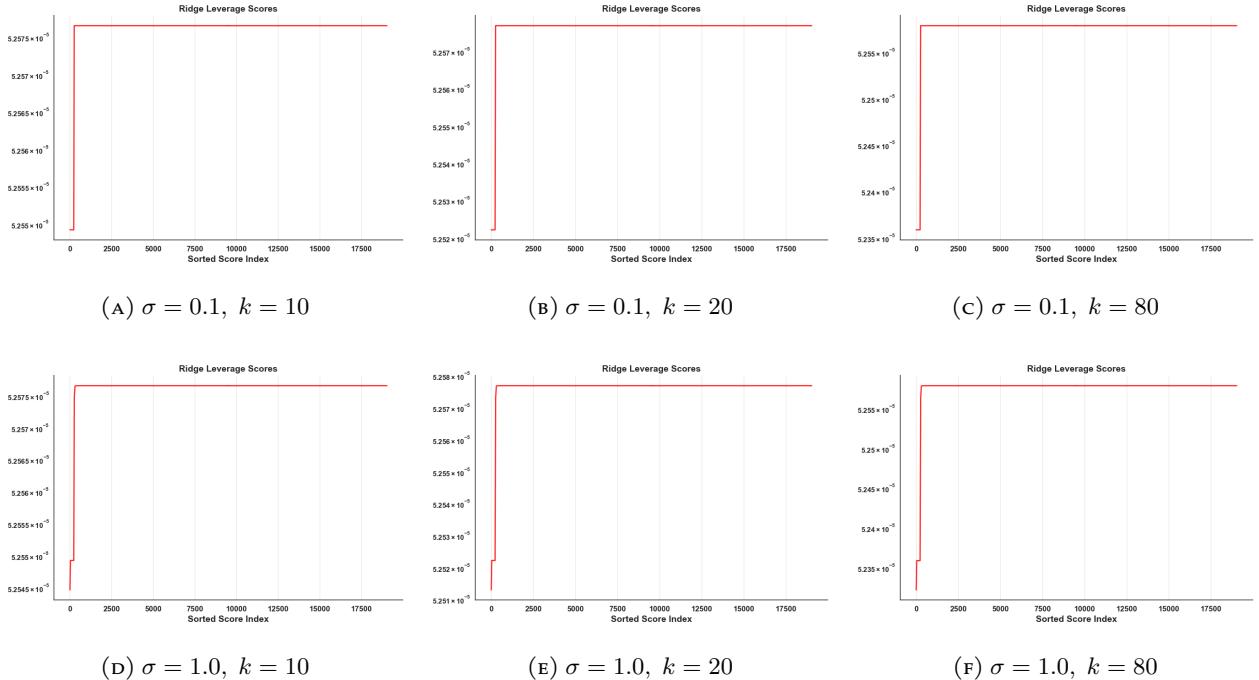


FIGURE 21. Ridge Leverage scores for the Magic data set.

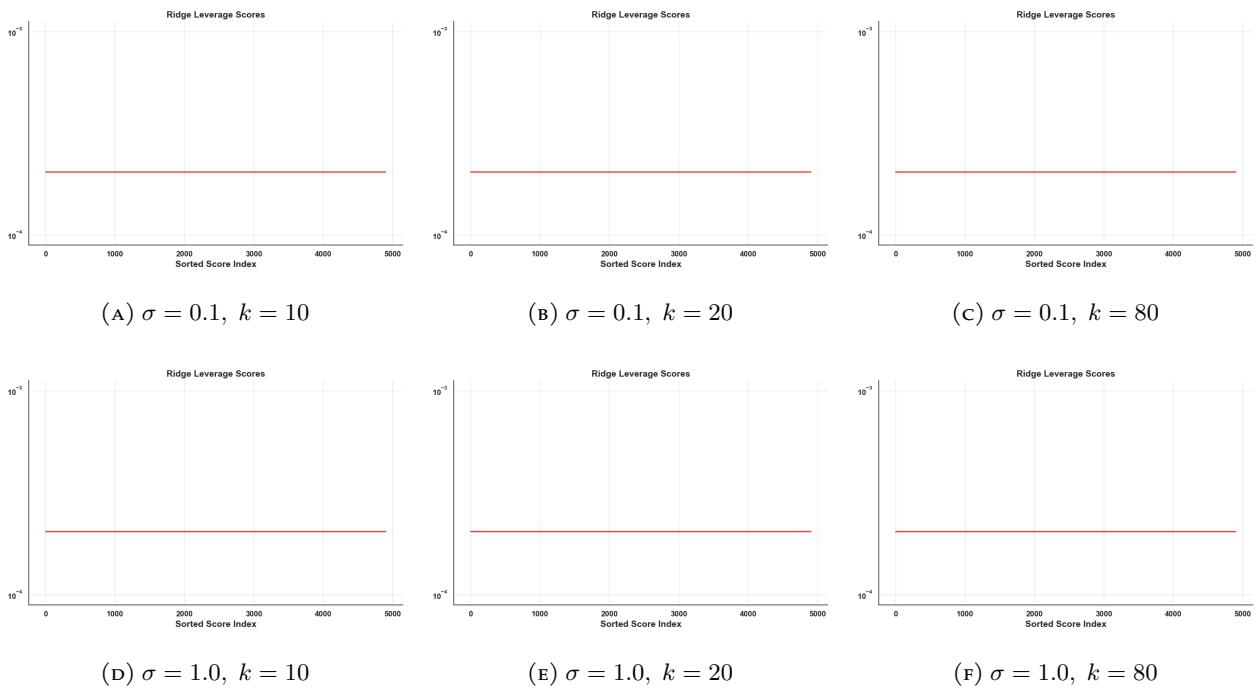


FIGURE 22. Ridge Leverage scores for the stock market data set.

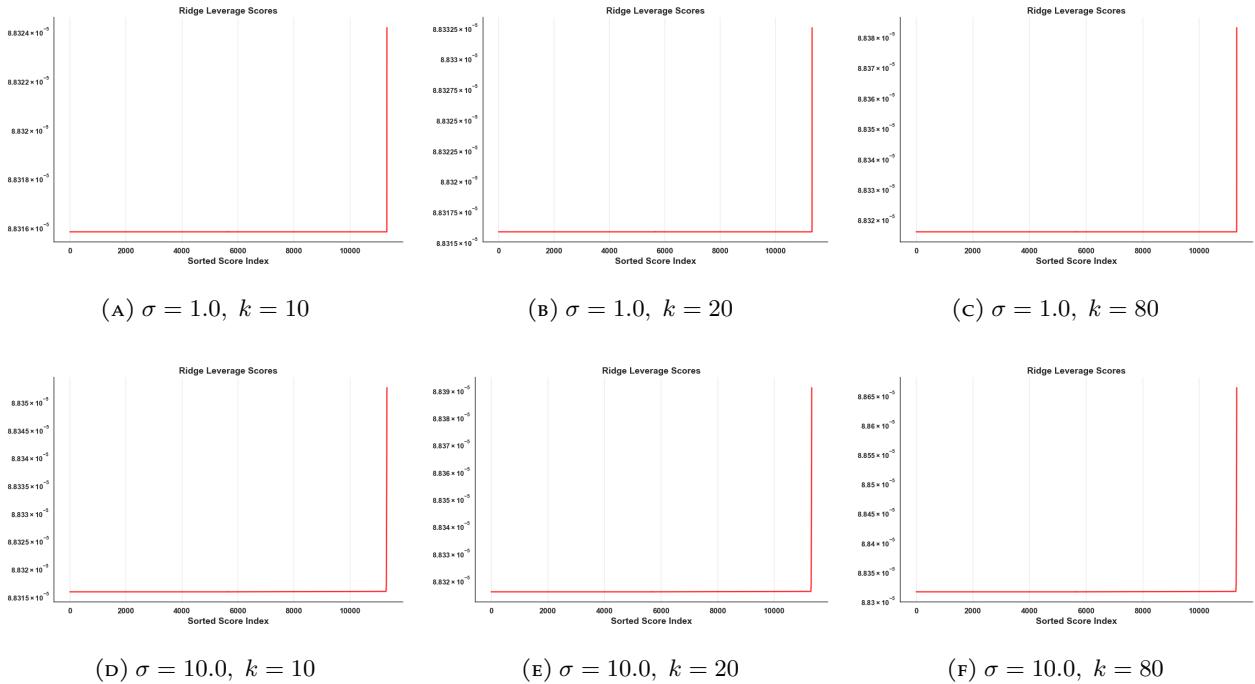


FIGURE 23. Ridge Leverage scores for the temperature data set.

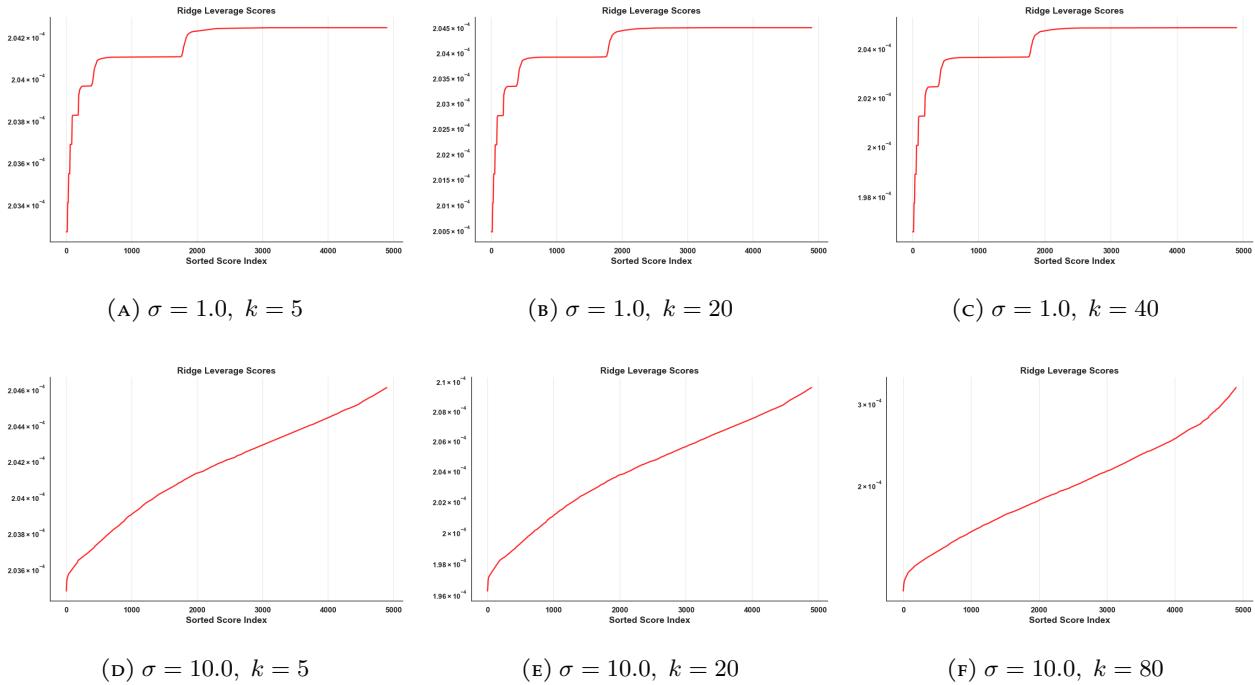


FIGURE 24. Ridge Leverage scores for the wine data set.

A.4. Nystrom Errors.

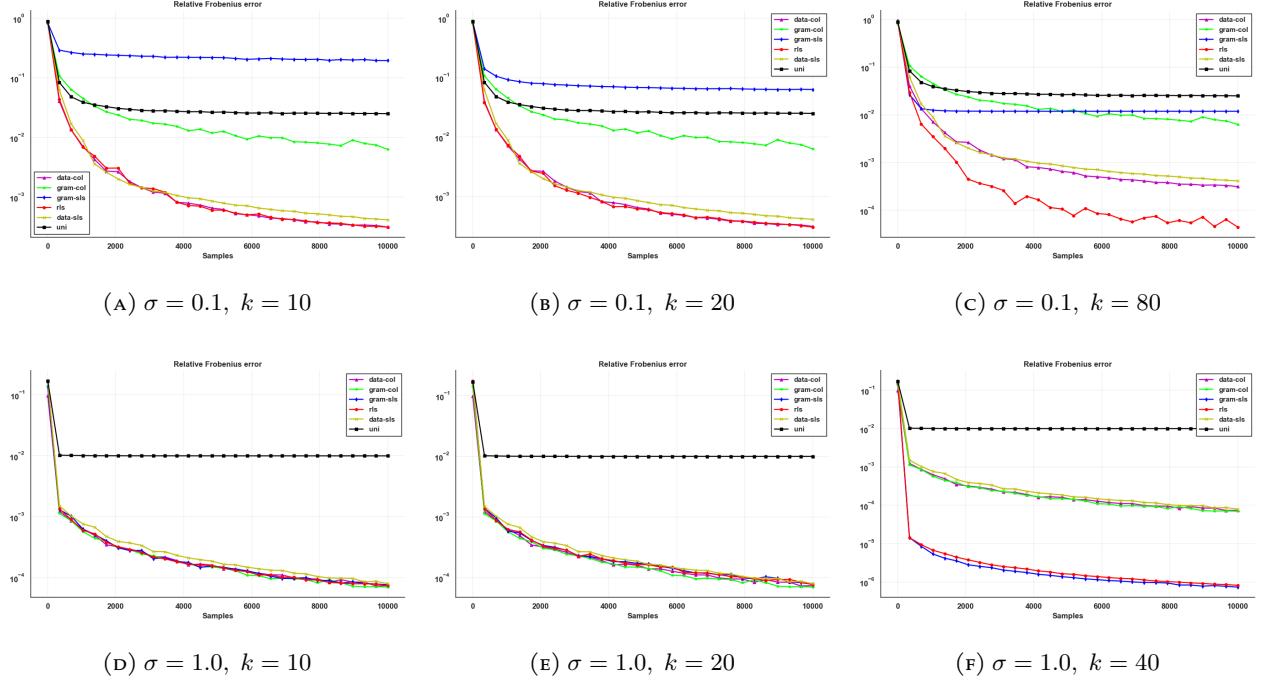


FIGURE 25. Nystrom Frobenius errors for the 3D-spatial network data set.

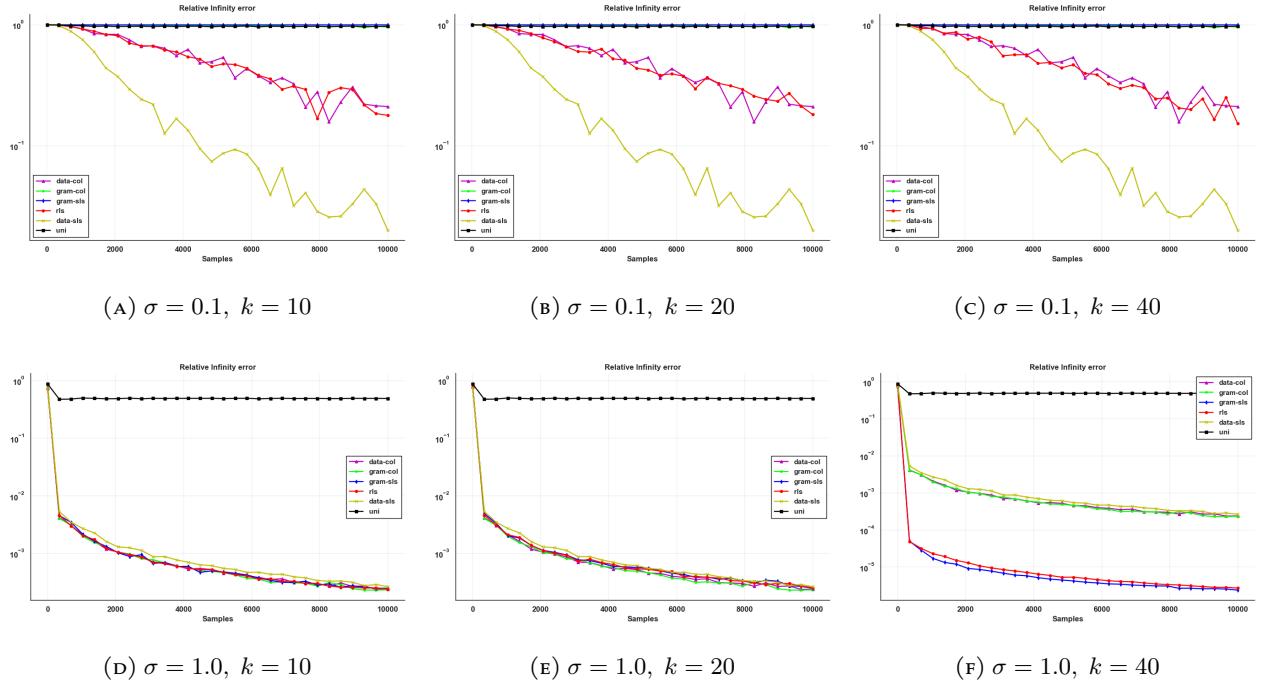


FIGURE 26. Nystrom absolute errors for the 3D-spatial network data set.

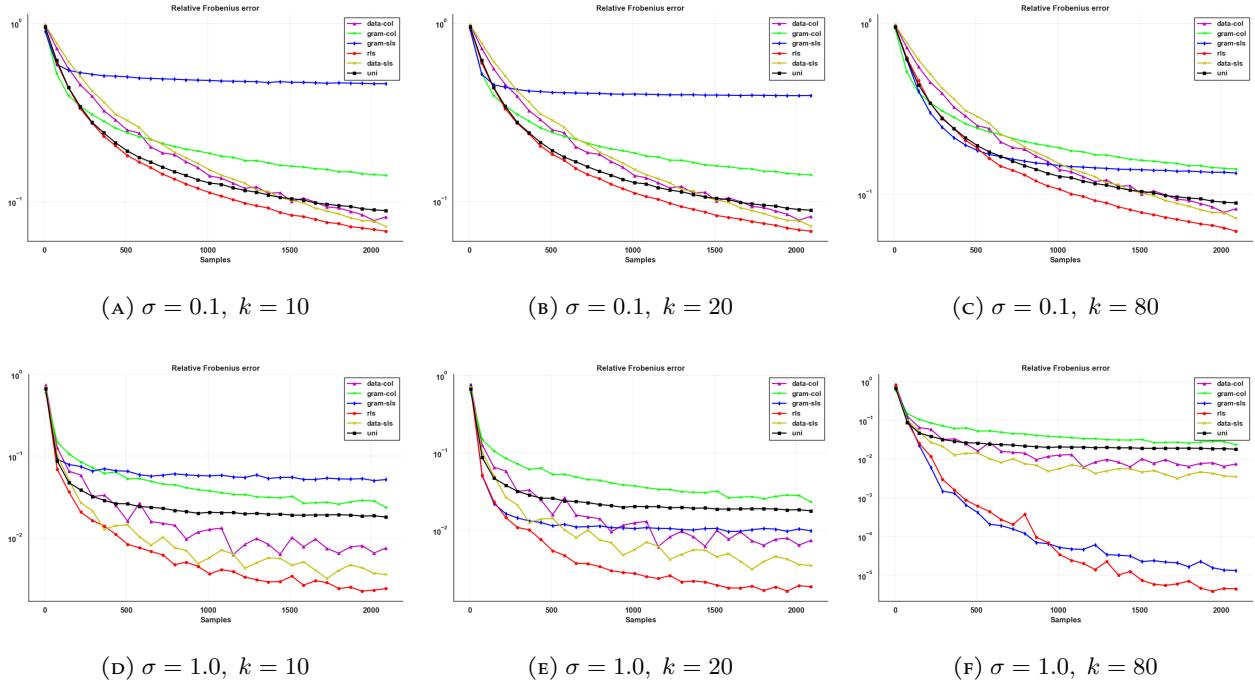


FIGURE 27. Nystrom Frobenius errors for the Abalone data set.

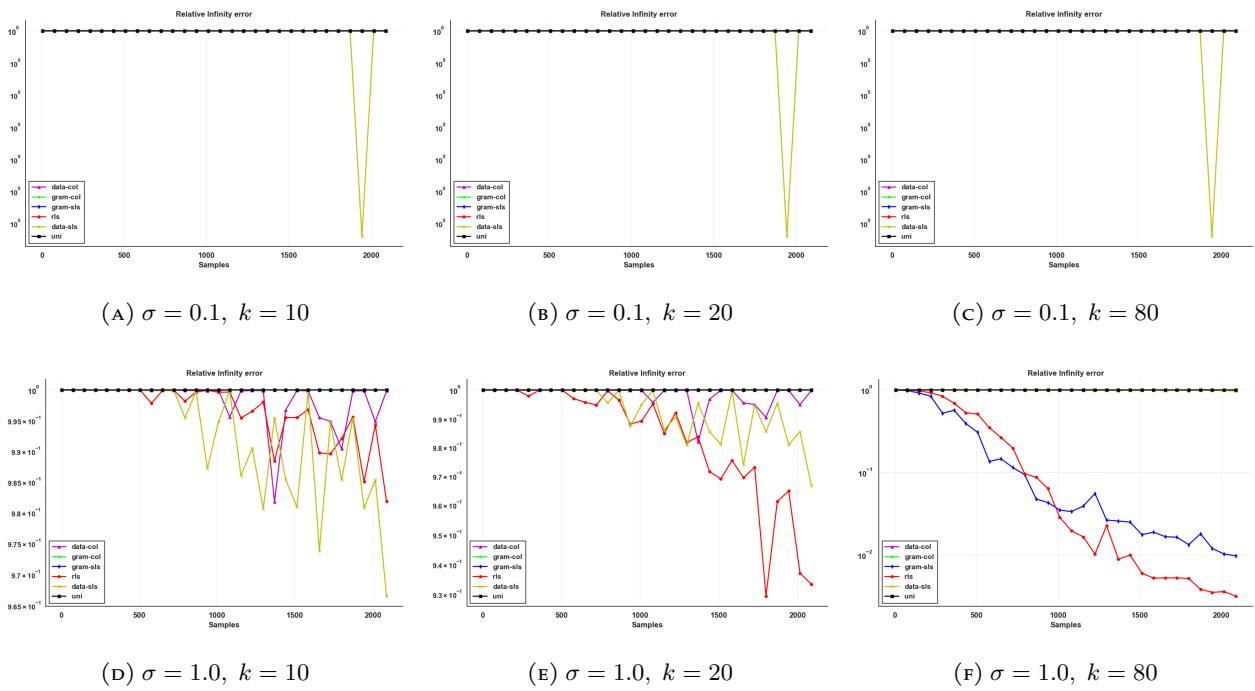


FIGURE 28. Nystrom absolute errors for the Magic data set.

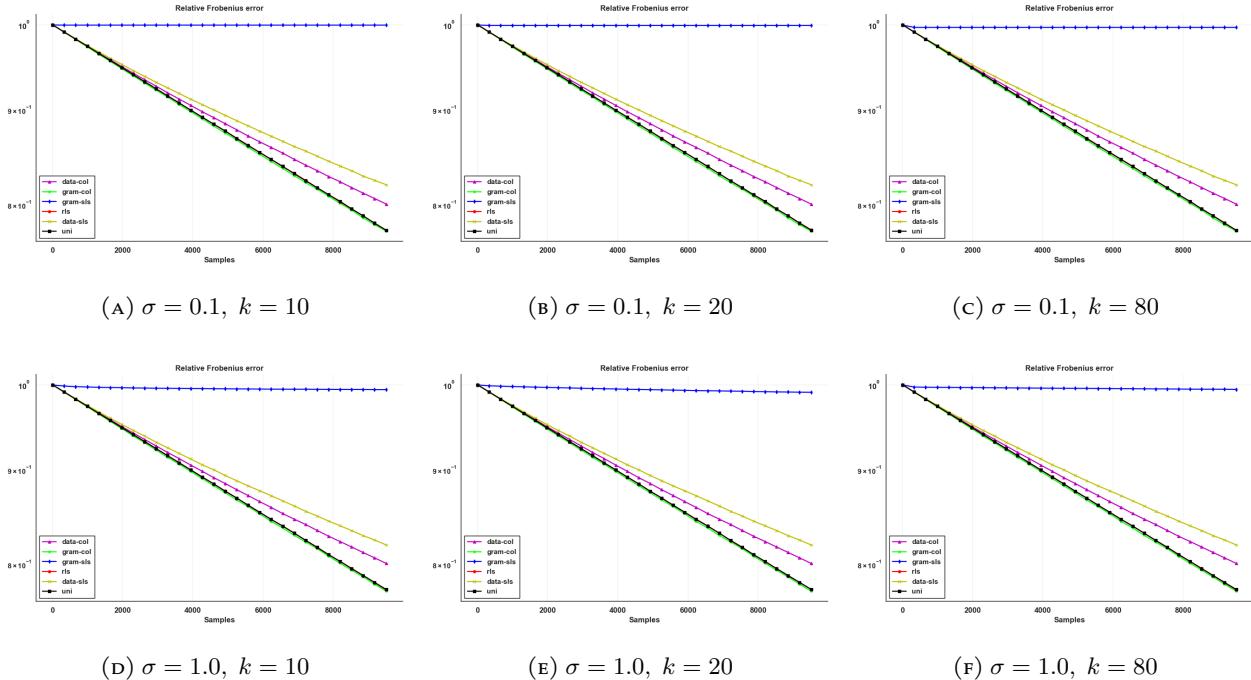


FIGURE 29. Nystrom Frobenius errors for the Magic data set.

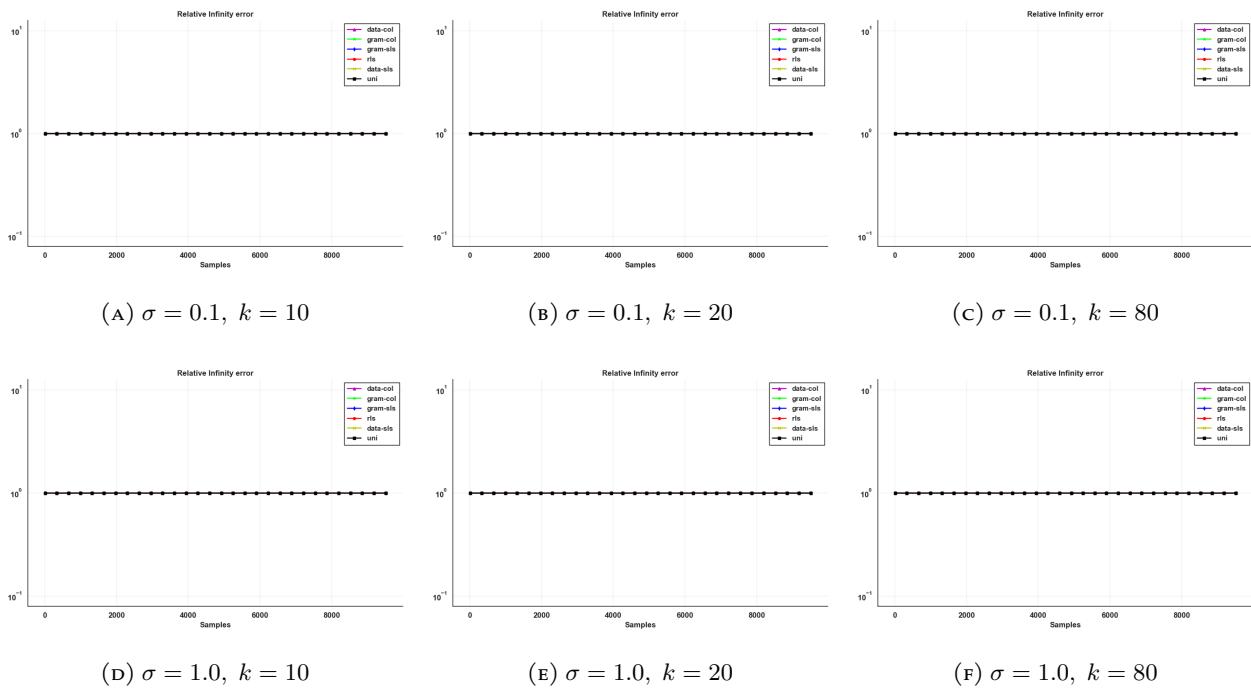


FIGURE 30. Nystrom absolute errors for the Magic data set.

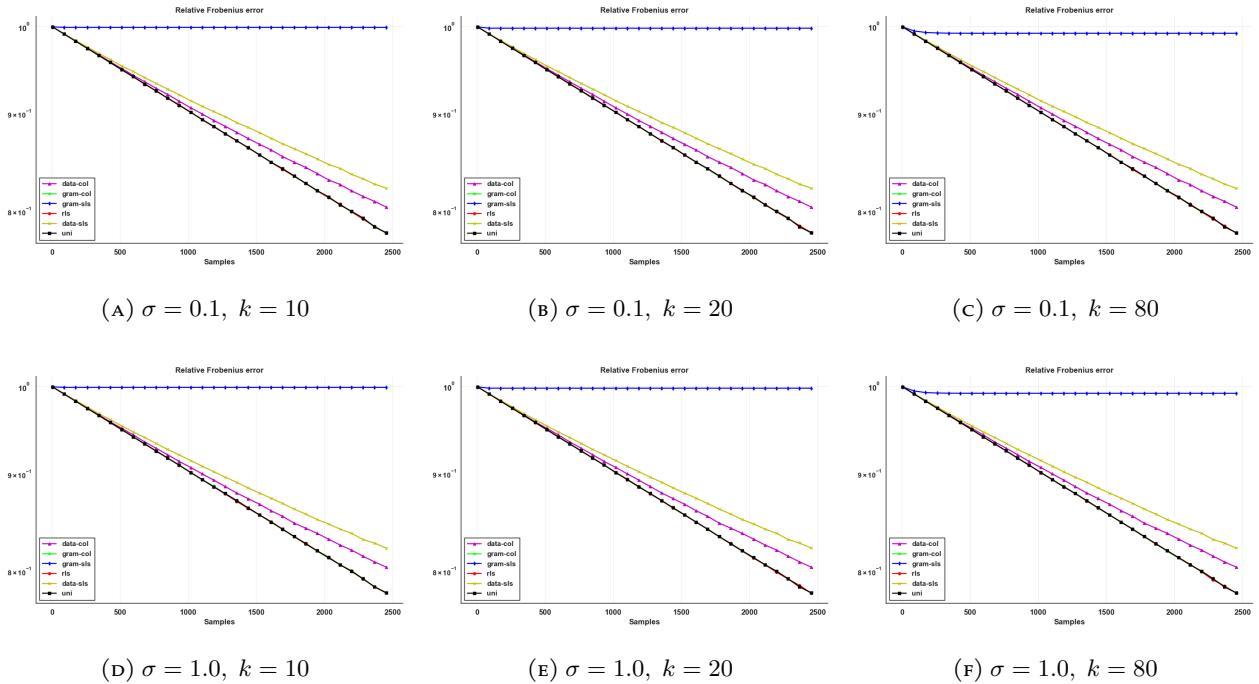


FIGURE 31. Nystrom Frobenius errors for the stock market data set.

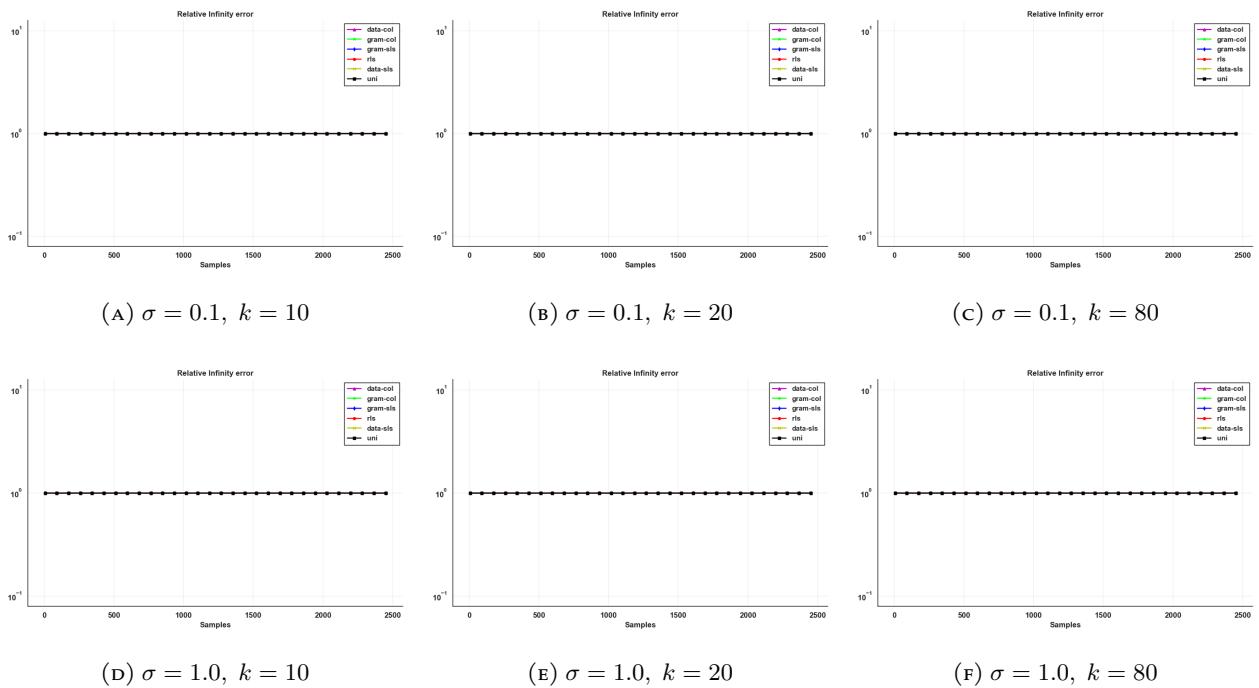


FIGURE 32. Nystrom absolute errors for the stock market data set.

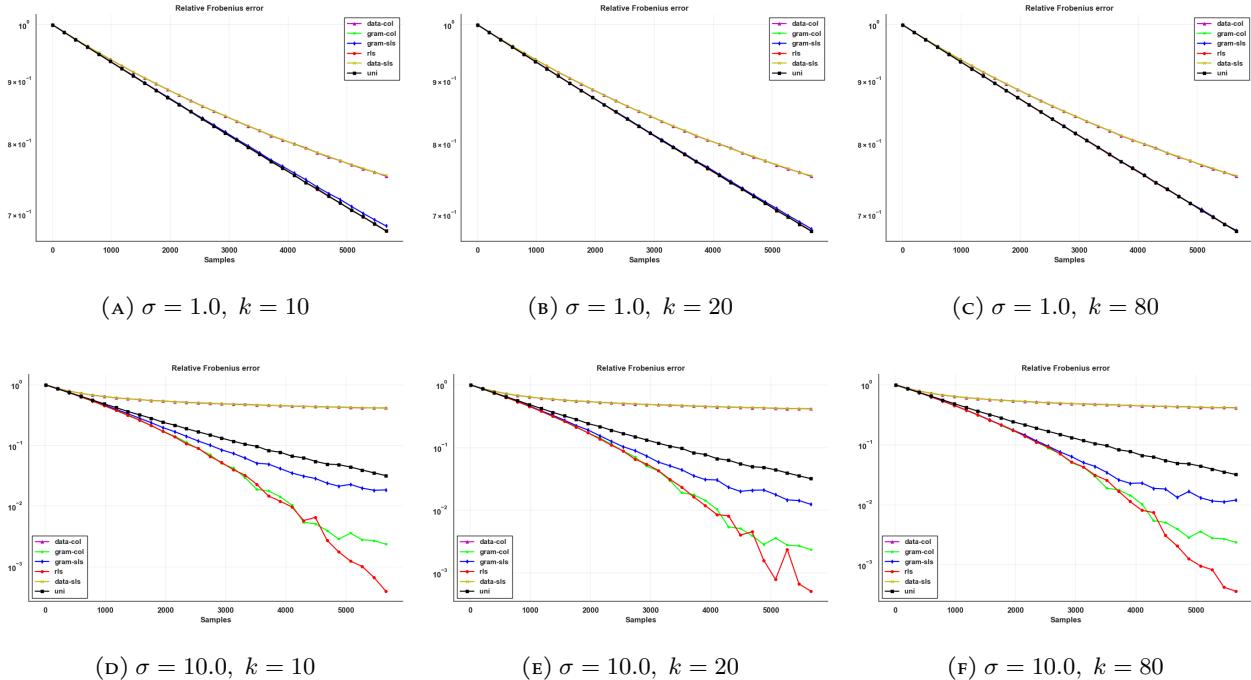


FIGURE 33. Nystrom Frobenius errors for the temperature data set.

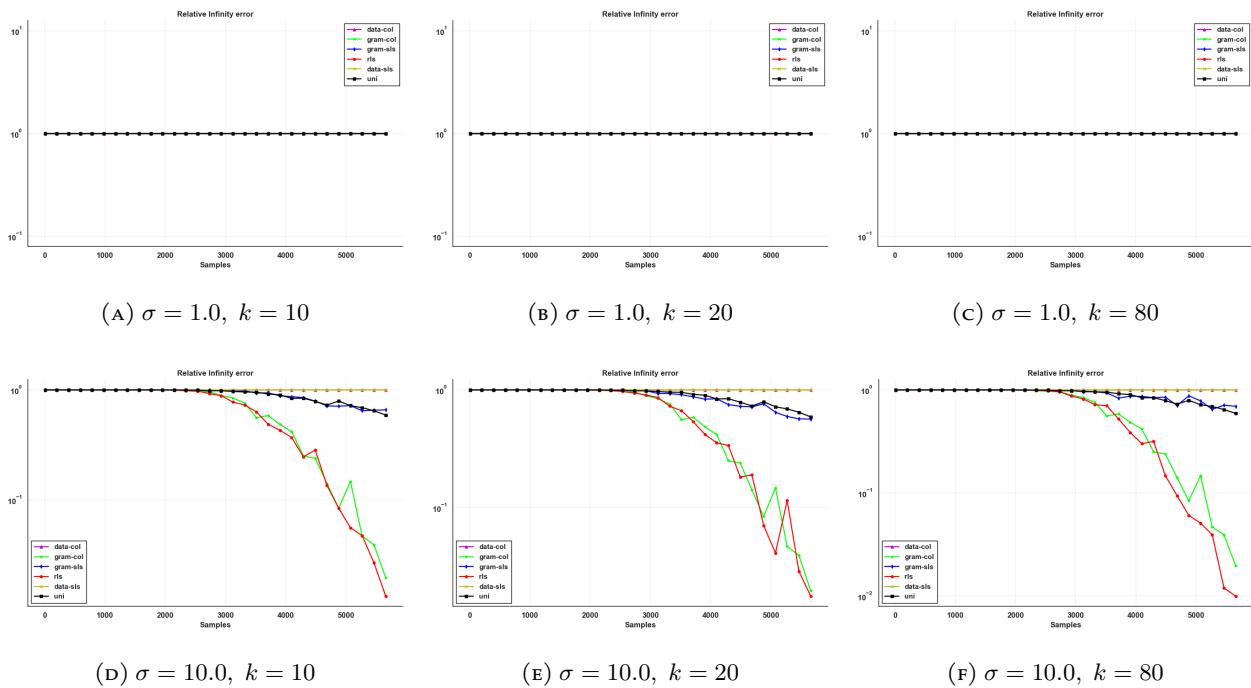


FIGURE 34. Nystrom absolute errors for the temperature data set.

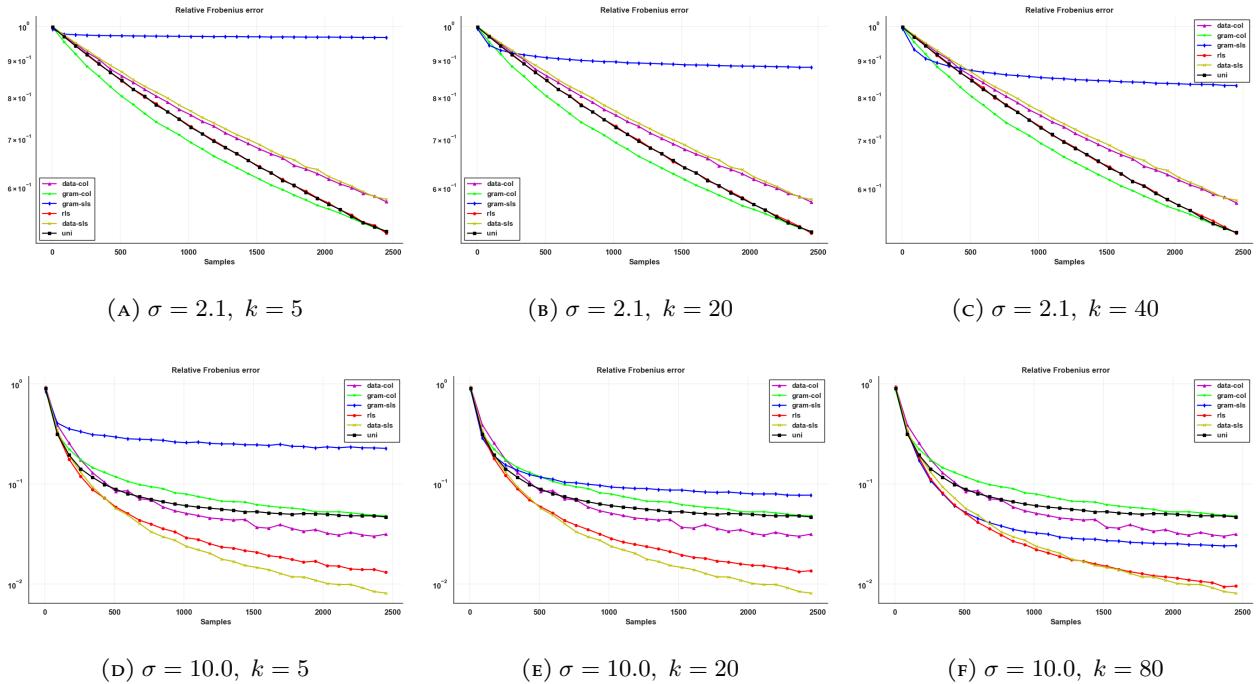


FIGURE 35. Nystrom Frobenius errors for the wine data set.

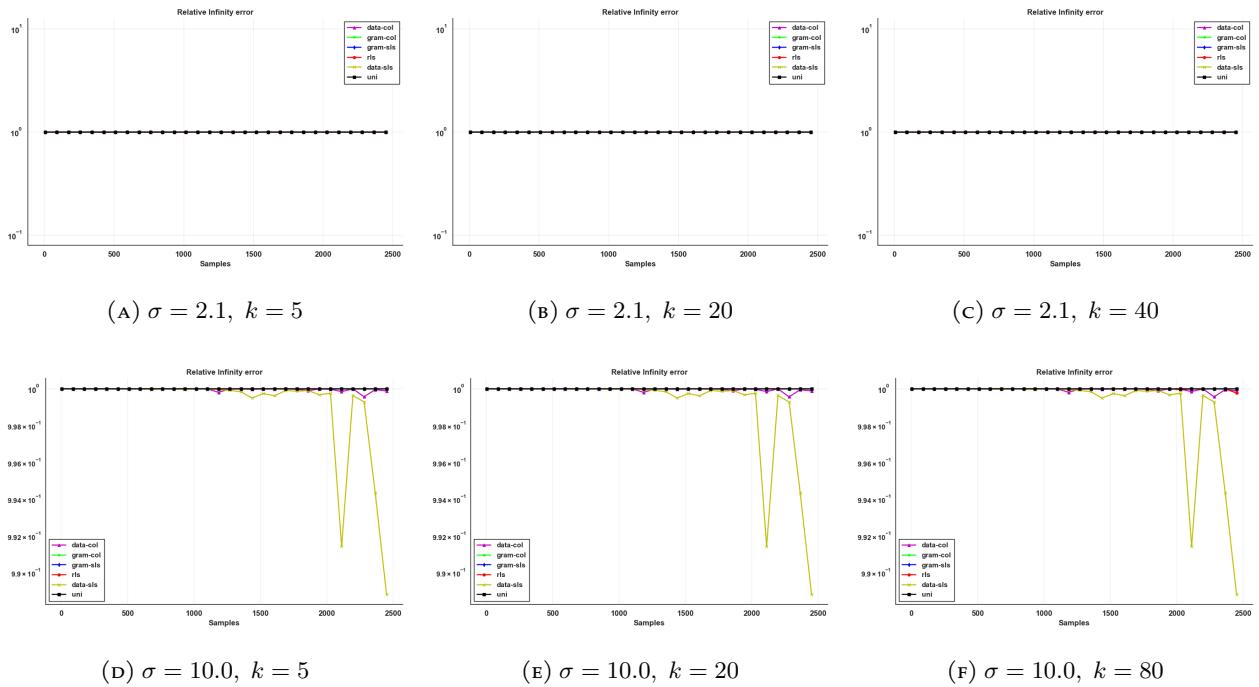


FIGURE 36. Nystrom absolute errors for the wine data set.